

Imperial College London

MENG INDIVIDUAL PROJECT

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

Systematic Benchmarking and Cost-Aware Profiling of Monte Carlo Workloads with Automatic Differentiation

Author:
Vivian Lopez

Supervisor:
Prof. William Knottenbelt

Second Marker:
Prof. Ben Glocker

June 11, 2026

Abstract

Monte Carlo (MC) simulation with automatic differentiation (AD) is widely used in quantitative finance and scientific computing, where both numerical estimates and parameter sensitivities must be computed efficiently. However, MC+AD performance is difficult to predict because runtime, memory behaviour, scalability, and numerical stability depend jointly on the stochastic workload, differentiation mode, execution backend, compiler/runtime behaviour, threading strategy, and hardware configuration. This makes it hard to choose appropriate software and cloud hardware for a given workload and deployment budget.

This project develops a reproducible benchmarking and profiling framework for MC+AD option-pricing workloads. The framework provides explicit workload specifications, a common engine interface, controlled warm-up and repeated measurement, deterministic seeding, numerical validation, structured result storage, and environment metadata capture. It evaluates European, local-volatility, and Asian option workloads across Python/NumPy, JAX, C++, and Rust implementations on local and Google Cloud CPU platforms. AD is evaluated primarily through JAX, while C++ and Rust provide compiled no-AD throughput baselines. The evaluation shows that no single backend or hardware configuration is uniformly optimal: backend rankings change with workload structure, scale, differentiation requirements, and cost constraints.

The project also implements a workload-aware cloud profiling methodology using short probe runs and successive-halving style promotion to reduce expensive full-scale benchmarking. SHA recovered the Pareto frontier in the main **n2-standard-4** case while halving full-scale evaluations, but cross-instance robustness was mixed, motivating guarded promotion checks rather than a claim of universal profiler reliability. Overall, the work provides a practical framework and empirical methodology for analysing MC+AD systems, and points towards future extensions in broader AD-system coverage, GPU evaluation, and more uncertainty-aware cloud profiling.

Acknowledgements

I would like to thank my supervisor, William Knottenbelt, for his helpful guidance and advice during the project. His feedback helped clarify the direction of the work and identify the areas where the project could be strengthened.

I am also grateful to the HSBC Quant team for their support and industrial perspective. In particular, I would like to thank Zouhair Rajehi for acting as the main point of contact and for providing useful background, technical context, and pointers for the project. I would also like to thank Renaud X. Delloye for his input during the meetings, and Ouejdane Ghodhbane for coordinating the meetings and providing a welcoming point of contact with the team.

Finally, I would like to thank the Department of Computing at Imperial College London for the opportunity to undertake this project.

Contents

1	Introduction	5
1.1	Motivation	5
1.2	Problem Statement	5
1.3	Research Questions	6
1.4	Contributions	6
1.5	Report Structure	6
2	Background and Related Work	8
2.1	Monte Carlo Simulation	8
2.2	Automatic Differentiation	8
2.3	Monte Carlo Greeks in Finance	9
2.4	Monte Carlo + AD Challenges	9
2.5	Existing Frameworks	10
2.6	Hardware Considerations	10
2.7	Benchmarking Methodology and Reproducibility	10
2.8	Compiler Techniques from Machine Learning Systems	10
2.9	Benchmark Reporting Standards	11
2.10	Budget-Aware Benchmark Selection	11
2.11	Summary of Gaps	11
3	System Design	12
3.1	Design Goals	12
3.2	High-Level Architecture	13
3.3	Layer 1 — Workload Configuration	13
3.3.1	Workload 1 — European Option	14
3.3.2	Workload 2 — European Local-Volatility Option	14
3.4	Layer 2 — Engine Interface	15
3.5	Layer 3 — Benchmark Runner	16
3.6	Layer 4 — Storage	16
3.7	Layer 5 — REST API	17
3.8	Layer 6 — Frontend	18
3.9	Cloud Cost-Performance Design	18
3.10	Budget-Aware Profiler Design	18
3.10.1	Probe-Based Baseline	18
3.10.2	Successive Halving Profiler	19
3.11	Design Decisions and Trade-offs	19
3.12	Summary	20
4	Implementation	21
4.1	Overview	21
4.2	Core Modules	21
4.2.1	Workload Registry	21
4.2.2	Monte Carlo Engine Interface	22
4.2.3	NumPy Engine	22
4.2.4	JAX Engine	22
4.2.5	C++ Engine	23
4.2.6	Rust Engine	23

4.2.7	CUDA Engine	24
4.2.8	Benchmark Runner	24
4.2.9	Metrics Collection	25
4.2.10	Storage Layer	25
4.2.11	REST API	25
4.2.12	Frontend / Dashboard	25
4.3	Supported Workloads	26
4.4	AD Integration	26
4.5	GPU Support	27
4.6	Challenges	27
4.7	Cloud Metadata and Cost Support	28
4.8	Profiler Implementation	28
4.9	Testing	28
4.10	Dependencies	29
4.11	Summary	29
5	Experimental Setup	30
5.1	Overview	30
5.2	Connection to the Research Questions	30
5.3	Execution Environment	30
5.4	Workloads	31
5.4.1	European Option under Geometric Brownian Motion	31
5.4.2	European Option under Parametric Local Volatility	31
5.4.3	Asian Arithmetic-Average Option	32
5.5	Engines and Software Stacks	32
5.6	Experimental Variables	32
5.6.1	Independent Variables	32
5.6.2	Dependent Variables	33
5.6.3	Controlled Variables	33
5.7	Metrics	34
5.8	Measurement Methodology	34
5.8.1	Warmup and JIT Compilation	34
5.8.2	Cloud Cost Model	35
5.8.3	Repeated Measurements	35
5.8.4	Randomness and Correctness	35
5.8.5	Fairness Controls	35
5.8.6	Data Exclusion Policy	35
5.9	Experiment Scripts	36
5.10	Individual Experiment Designs	36
5.11	Final Cloud and Profiler Experiments	36
5.12	Execution Order and Analysis Queries	38
5.13	Reproducibility	38
5.14	Summary	38
6	Results and Discussion	39
6.1	Overview	39
6.2	European GBM Validation	40
6.3	Local-Volatility Smoke Test	40
6.4	Local-Volatility Engine Scalability	41
6.5	Local-Volatility AD Scaling with Path Count	41
6.6	Local-Volatility Scaling with Time Steps	42
6.7	Local-Volatility AD Scaling with Time Steps	42
6.8	Memory Pressure and Practical Resource Limits	43
6.9	Thread Scalability on Local Volatility	43
6.9.1	Strong Scaling	43
6.9.2	Weak Scaling	44
6.10	Surface-Shape Sensitivity	44
6.11	Timing Stability	44
6.12	Stress Limit	45

6.13	Final Cloud Cost-Performance and Profiler Results	45
6.13.1	Correctness and AD Overhead in the Final Cloud Run	45
6.13.2	Cloud Cost-Performance	46
6.13.3	Pareto Frontier	46
6.13.4	Old Profiler versus SHA	48
6.13.5	Guarded SHA Local-Volatility Profiler	48
6.13.6	SHA Progression in the Earlier Multi-Workload Grid	49
6.14	Cross-Question Discussion	50
6.14.1	Software Stack and Engine Performance	50
6.14.2	Programming-Language, Runtime and AD-Support Trade-offs	50
6.14.3	Parallelism and Bottlenecks	50
6.14.4	Compiler and AD Runtime Techniques	50
6.14.5	Integrated Scope of the Completed Evaluation	51
6.15	Limitations and Threats to Validity	51
6.15.1	Hardware and Cloud Scope	51
6.15.2	Local-Volatility Correctness	51
6.15.3	Random Number Generation as a Confounder	51
6.15.4	Warm versus Cold JAX Execution	52
6.15.5	Cloud Cost Scope	52
6.15.6	Memory Measurement	52
6.15.7	Thread Control	52
6.16	Summary	52
7	Conclusion	53
7.1	Summary of Work	53
7.2	Key Findings	53
7.3	Answering Research Questions	54
7.4	Future Work	54
7.5	Final Remarks	55
	Declarations	56

Chapter 1

Introduction

1.1 Motivation

Monte Carlo (MC) simulation is a foundational technique in scientific computing and quantitative finance, used to estimate expectations of stochastic systems where closed-form solutions are unavailable. In financial applications, MC methods are widely used for pricing complex derivatives and evaluating risk under uncertainty. However, modern use cases increasingly require not only accurate estimates, but also sensitivities of these estimates with respect to model parameters, such as Greeks in option pricing.

Automatic differentiation (AD) provides a principled and efficient way to compute such sensitivities by differentiating the numerical program itself. Unlike finite difference methods, AD produces derivatives that are accurate up to floating-point precision and integrates naturally with simulation code. The combination of MC and AD (MC+AD) therefore enables flexible, gradient-based analysis of stochastic systems.

Despite its advantages, MC+AD introduces significant computational challenges. These workloads are typically compute-intensive, memory-sensitive, and highly dependent on implementation details such as hardware architecture, programming language, numerical libraries, and differentiation mode. As high-performance computing (HPC) systems become increasingly heterogeneous—spanning CPUs, GPUs, and cloud-based environments—understanding how MC+AD workloads behave across these systems becomes critical.

This project is motivated by the need to systematically evaluate and optimise MC+AD workloads in modern computing environments, with a focus on performance, scalability, and cost-efficiency.

1.2 Problem Statement

While Monte Carlo simulation, automatic differentiation, and high-performance computing have each been studied extensively, there is a lack of unified frameworks and methodologies for evaluating their combined behaviour. Existing work often focuses on isolated aspects, such as MC algorithms, AD systems, or hardware benchmarking, without considering their interactions in a controlled and reproducible manner.

In particular, several gaps can be identified. First, there is no standardised benchmarking framework that supports MC workloads, explicit AD-capable backends, and comparable non-AD compiled baselines across multiple programming languages and hardware architectures. Second, most AD benchmarking studies are conducted in the context of machine learning, leaving stochastic simulation workloads comparatively underexplored. Third, performance comparisons across systems are often difficult to interpret due to inconsistent experimental setups, lack of reproducibility, and insufficient reporting of configuration details.

Furthermore, compiler techniques developed in the machine-learning ecosystem, such as staged execution, graph compilation, and operator fusion, have not been systematically evaluated in the context of MC+AD workloads. Finally, there is limited empirical guidance on selecting appropriate cloud hardware configurations for these workloads, particularly when considering cost-performance trade-offs.

As a result, practitioners lack clear, evidence-based insights into how to design, implement, and deploy efficient MC+AD systems.

1.3 Research Questions

This project addresses the following research questions. The same numbering is used in the experimental design and conclusion so that each result can be traced back to the problem framing.

- RQ1.** How do different software stacks and execution backends affect the runtime, throughput, and memory behaviour of Monte Carlo workloads?
- RQ2.** What are the performance and memory trade-offs of automatic differentiation for representative Monte Carlo option-pricing workloads?
- RQ3.** How do workload structure and scale, including path count, time-step count, and payoff type, affect backend rankings and bottlenecks?
- RQ4.** Which Google Cloud CPU configurations provide the best marginal cost-performance for the evaluated MC+AD workloads?
- RQ5.** Can short probe runs and successive-halving style profiling reduce the cost of cloud benchmarking while still recovering near-Pareto-optimal configurations?
- RQ6.** What are the practical limitations, threats to validity, and generalisation limits of this benchmarking methodology?

1.4 Contributions

This project makes the following contributions:

- The design and implementation of a modular, extensible benchmarking framework for Monte Carlo engines with automatic differentiation support.
- A systematic methodology for evaluating Monte Carlo workloads across different hardware architectures, software stacks, and programming languages, with JAX providing the main implemented AD case study.
- An experimental study comparing performance, scalability, and memory behaviour under varying configurations, including JAX forward/reverse AD modes and native compiled no-AD baselines.
- An analysis of cost-performance trade-offs for MC+AD workloads in cloud computing environments, including measured Google Cloud runs and per-run cost estimates.
- A budget-aware profiling methodology based on short probe runs and successive halving, evaluated against a full-grid cloud baseline.
- Insights into the applicability and limits of JAX/XLA-style compiler/runtime techniques for stochastic simulation workloads.

The framework and methodology developed in this work provide a foundation for reproducible and extensible benchmarking of MC+AD systems. The final evaluation combines detailed local CPU experiments with cloud cost-performance and profiler-selection experiments.

The strongest quantitative conclusions are therefore scoped to the implemented CPU engines, the JAX AD path, and the Google Cloud CPU experiments evaluated in this report. GPU support and broader AD-system coverage are present as prototype infrastructure or future extensions, but they are not used as the basis for the final comparative claims.

1.5 Report Structure

The remainder of this report is organised as follows.

- Chapter 2 reviews the background and related work, including Monte Carlo simulation engines, automatic differentiation, benchmarking methodologies, and relevant techniques from high-performance computing and machine learning.
- Chapter 3 describes the system design of the benchmarking framework, including workload definitions, execution model, data collection pipeline, cloud metadata model, and budget-aware profiling design.

- Chapter 4 describes the concrete implementation of the framework, including the engine modules, benchmark runner, storage layer, API, dashboard, cloud-cost metadata, and profiler driver.
- Chapter 5 outlines the experimental methodology, covering validation, performance benchmarking, scalability analysis, cloud experiments, and profiler-versus-grid comparisons.
- Chapter 6 presents and analyses the experimental results, focusing on performance trends, AD overheads, scaling limits, cloud cost-performance, Pareto frontiers, and profiler quality.
- Chapter 7 concludes the report, summarising key findings and discussing potential directions for future work.

Chapter 2

Background and Related Work

2.1 Monte Carlo Simulation

Monte Carlo (MC) simulation is a widely used technique for estimating expectations of stochastic systems through repeated random sampling [1]. It is particularly valuable in settings where analytical solutions are unavailable or intractable, such as the pricing of financial derivatives under complex dynamics. In such applications, MC methods approximate quantities of interest by simulating a large number of independent sample paths and averaging the resulting payoffs.

A key property of MC simulation is that its statistical error decreases at a rate proportional to $O(N^{-1/2})$, where N is the number of samples. As a result, achieving high accuracy typically requires a large number of simulations, making computational efficiency a primary concern. At the same time, MC workloads are often described as “embarrassingly parallel”, since individual simulation paths can be evaluated independently. This makes them well-suited to modern parallel hardware, including multi-core CPUs and GPUs.

In practice, Monte Carlo *engines* extend basic simulations by providing configurable, reusable implementations that support parameter variation, reproducibility through controlled random number generation, and repeated execution for statistical aggregation. Benchmarking such engines therefore requires not only measuring raw execution time, but also understanding how performance scales with workload size, parallelism, and implementation choices.

A useful distinction for this project is between a Monte Carlo engine and a benchmarking framework. An engine implements the numerical workload itself: for example, a European option-pricing engine may accept a strike, volatility, interest rate, maturity, path count, and random seed, then return a price and possibly sensitivities. A framework sits around such engines. It is responsible for constructing repeated runs, changing parameters systematically, measuring runtime and memory behaviour, recording hardware metadata, and producing analysis artefacts. This distinction is important because many performance claims about Monte Carlo code are not caused by the payoff formula alone, but by the surrounding machinery: random-number generation, warm-up policy, thread configuration, storage, and postprocessing.

2.2 Automatic Differentiation

Automatic differentiation (AD) is a technique for computing derivatives of numerical programs by systematically applying the chain rule to the sequence of elementary operations that define the computation [2, 3]. Unlike finite difference methods, AD produces derivatives accurate up to floating-point precision while avoiding numerical instability.

Two primary modes of AD are commonly used. Forward-mode AD propagates derivatives alongside values during the forward execution of a program and is typically efficient when the number of input parameters is small relative to the number of outputs. In contrast, reverse-mode AD records intermediate computations during a forward pass and then propagates sensitivities backwards, making it more efficient when a small number of outputs depend on many inputs.

Modern AD systems, such as those integrated into scientific computing and machine learning frameworks, enable gradients to be computed directly from program implementations. JAX is one example of this style of system, combining NumPy-like array programming with composable transformations such as JIT compilation and automatic differentiation [4]. This is particularly

useful in simulation-based contexts, where sensitivities (e.g., financial “Greeks”) can be obtained without manual derivation.

However, the choice between forward- and reverse-mode AD introduces important trade-offs in computational cost and memory usage. Forward-mode typically incurs moderate overhead with minimal additional memory requirements, while reverse-mode can significantly increase memory usage due to the need to store intermediate states. These trade-offs become especially important in large-scale simulations.

Several AD systems are relevant when considering how MC+AD support might be extended beyond the JAX implementation evaluated in this project. Enzyme differentiates LLVM intermediate representation, making it attractive for C and C++ style numerical kernels [5]. Zygote provides source-to-source AD in the Julia ecosystem [6], while DiffTaichi demonstrates differentiable simulation on accelerator-oriented programming abstractions [7]. These systems illustrate that AD is not a single implementation technique. The design choice for this project was therefore to treat AD mode as part of the benchmark configuration, rather than assuming that all engines expose identical derivative capability.

2.3 Monte Carlo Greeks in Finance

In quantitative finance, sensitivities such as Delta, Vega and Rho are often as important as the option price itself. Standard Monte Carlo Greek estimators include finite differences, pathwise derivative methods, and likelihood-ratio or score-function methods [1, 8]. Finite differences are simple but introduce step-size and noise trade-offs; common random numbers reduce variance but do not remove the bias-variance choice. Pathwise and likelihood-ratio estimators avoid some of these problems under appropriate smoothness or density assumptions.

Adjoint algorithmic differentiation extends this idea computationally: once a pricing computation is written as a program, reverse accumulation can produce many sensitivities at a cost that is often a small multiple of the primal valuation [9]. This project does not propose a new Greek estimator. Existing work studies Monte Carlo Greeks and AD/AAD methods, while this project studies the systems-level benchmarking, reproducibility and cost-performance behaviour of implemented MC+AD workloads across backends and cloud CPU configurations.

2.4 Monte Carlo + AD Challenges

Combining Monte Carlo simulation with automatic differentiation enables flexible and general sensitivity analysis, but introduces several non-trivial computational challenges.

First, the computational cost of AD compounds with the cost of simulation. In MC workloads, each sample path may consist of many time steps and numerical operations. Applying AD to such computations increases the amount of work performed per path, with the overhead depending on the differentiation mode and implementation. As a result, throughput (e.g., paths per second) can be significantly reduced compared to non-differentiated simulations.

Second, memory usage becomes a critical concern, particularly for reverse-mode AD. Reverse-mode requires storing intermediate values during the forward execution in order to compute gradients during the backward pass. In long-horizon simulations with many time steps, this can lead to substantial memory growth, potentially exceeding cache capacity and degrading performance due to increased memory traffic. In extreme cases, memory limitations may prevent scaling to larger workloads altogether.

Third, the interaction between AD and parallelism is complex. While Monte Carlo simulation is naturally parallel across independent paths, AD introduces dependencies within each path that may limit opportunities for parallel execution. Furthermore, reverse-mode AD can introduce additional synchronisation points and data movement, particularly in GPU or distributed settings, reducing the effectiveness of parallelisation strategies.

Finally, the combination of stochastic computation and differentiation introduces challenges for reproducibility and numerical stability. Ensuring that differentiated simulations produce consistent and meaningful results requires careful control of random number generation, seeding, and execution order.

These challenges motivate the need for systematic evaluation of MC+AD workloads, where both computational performance and memory behaviour are measured under controlled and reproducible

conditions.

2.5 Existing Frameworks

Several frameworks exist for Monte Carlo simulation and high-performance numerical computing, though few target MC+AD benchmarking across heterogeneous backends. Carlo.jl is a relevant example because it separates the Monte Carlo algorithm from scheduling, checkpointing, storage and postprocessing, and it emphasises reproducible execution [10]. Those ideas influenced this project: workload definitions, engine execution, measurement policy and storage are separated so that benchmark rows remain auditable when engines are added or replaced.

However, Carlo.jl is intentionally scoped to a single language ecosystem and does not directly address cross-language benchmarking, AD support across backends, or cloud cost-performance. More broadly, scientific and machine learning benchmarks often focus on linear algebra kernels or neural network training rather than the interaction between stochastic simulation, differentiation and parallel execution. This motivates a framework that treats Monte Carlo simulation, AD-capable execution, reproducibility metadata and cloud cost as first-class benchmarking concerns.

2.6 Hardware Considerations

MC and MC+AD performance depends on arithmetic throughput, memory bandwidth, cache behaviour, vectorisation, multi-threading and runtime overheads. GPUs can provide high throughput for large data-parallel simulations, but kernel launch, memory transfer, control-flow divergence and AD memory requirements complicate the comparison. In cloud environments, raw speed must also be weighed against instance price and utilisation. The evaluation therefore reports both runtime and cost-performance rather than assuming that the fastest machine is the best deployment choice.

2.7 Benchmarking Methodology and Reproducibility

Benchmarking numerical workloads is difficult because small methodological differences can change the result. Warm-up effects, JIT compilation, thread pool initialisation, random-number generation, cache state, and background system load can all distort a single timing. This is particularly important for MC+AD workloads, where the cost per path may vary with the number of time steps, the differentiation mode, and the memory layout used by the engine.

For low-level CPU studies, benchmark suites such as PolyBench/C are useful because they isolate common loop nests and memory-access patterns under controlled conditions [11]. MC+AD workloads are higher level than these kernels, but the same lesson applies: the benchmark must make the execution conditions explicit. This project therefore uses warm-up runs, repeated timed measurements, deterministic seeds where supported, and database records that capture both workload parameters and environment metadata.

Hardware-counter and profiler tools provide another perspective on performance. LIKWID exposes CPU topology and performance counters in a lightweight workflow [12], while Intel VTune is widely used for deeper CPU performance analysis [13]. For XLA-backed systems, OpenXLA’s profiling documentation illustrates how modern compiler runtimes surface kernel-level and memory behaviour [14]. These tools were not used as primary evidence in the final results, but they inform the framework design: runtime alone is not always enough to explain why an engine scales or fails to scale, so metadata capture and profiler integration are natural extensions of the same benchmark lifecycle.

2.8 Compiler Techniques from Machine Learning Systems

Several compiler techniques associated with machine-learning systems are relevant to scientific simulation. Operator fusion and graph compilation can reduce memory traffic and dispatch overhead, which matters for Monte Carlo workloads that repeatedly apply similar update kernels. MLIR provides one example of an extensible compiler infrastructure for lowering computations to different hardware targets [15]. TVM demonstrates automated code generation and optimisation for tensor programs [16], and Ansoir extends this direction with hardware-aware search [17].

The relevance to this project is not that MC+AD workloads are identical to neural-network training, but that both domains benefit from expressing computation in a form that compilers can optimise. In the implementation chapter, this appears concretely in the use of JAX/XLA and `jax.lax.scan` for the local-volatility loop. Expressing the loop as a staged computation allows XLA to compile it as a structured loop rather than repeatedly dispatching Python operations. The evaluation therefore treats JAX/XLA as an AD-capable compiled runtime whose behaviour should be measured against NumPy and native compiled engines, rather than assumed to be faster or slower by design.

Mixed precision is another technique that has been successful in machine learning workloads [18]. For Monte Carlo simulation, the question is more delicate because reducing precision can affect both statistical estimates and pathwise derivatives. This project records precision and backend metadata where available, but the final evaluation focuses on double-precision CPU-oriented runs. The broader point is that benchmark configuration should expose such choices explicitly instead of hiding them as engine-specific optimisations.

2.9 Benchmark Reporting Standards

Large-scale benchmarks such as DAWNBench and MLPerf emphasise transparent reporting of workload, hardware, software versions and measurement protocol [19, 20, 21]. This report applies the same discipline at project scale: tables identify workload, engine, AD mode, path count and instance type, while the experiments chapter states warm-up, repetition, validation and exclusion policies.

2.10 Budget-Aware Benchmark Selection

Full-grid benchmarking is attractive because it gives a direct comparison of all candidate workloads, engines, AD modes, and hardware choices. However, it becomes expensive when each candidate must be evaluated at production-scale path counts on cloud instances. A practical benchmarking system therefore needs not only accurate measurement, but also a way to decide which configurations deserve the most expensive measurements.

This motivates the use of cheap probe runs. A probe run evaluates a candidate configuration at a smaller path count or reduced budget and uses that result as evidence about whether the configuration is likely to remain competitive at full scale. The central risk is selection error: a configuration that is slow at low budget may become competitive at high budget, or a noisy probe may give an incorrect ranking. Any profiler must therefore be evaluated against a full-grid baseline rather than assumed to be correct.

Successive halving is a natural fit for this setting. It begins by evaluating all candidates at a low budget, promotes a smaller active set to a larger budget, and repeats until only the strongest candidates remain for full evaluation. In the context of MC+AD benchmarking, the budget can be interpreted as the number of simulated paths. This connects the statistical structure of Monte Carlo experiments with the systems problem of cloud cost control: most configurations receive only cheap measurements, while promising configurations receive enough budget to support reliable final decisions.

2.11 Summary of Gaps

The literature motivates three gaps addressed by this project. First, existing Monte Carlo frameworks rarely combine reproducible MC execution, AD-capable workloads, cross-backend comparison and cloud metadata in one benchmark loop. Second, AD performance in stochastic simulation is less standardised than in machine-learning benchmarks, particularly for memory pressure and parallel execution. Third, practical deployment requires cost-performance and profiler selection evidence, not only raw runtime. This project addresses these gaps by building a modular benchmarking framework and using it to evaluate implemented MC+AD workloads across software stacks and cloud CPU configurations.

Chapter 3

System Design

3.1 Design Goals

The framework is designed around four guiding principles: *reproducibility*, *extensibility*, *fairness*, and *traceability*. Together these principles constrain almost every interface, data structure and naming convention in the codebase.

Reproducibility. Every benchmark run is fully described by a single serialisable workload configuration object. The `config_hash()` method serialises this object into canonically sorted JSON and computes a truncated (8-character) SHA-256 digest. The hash is stable across machines and Python versions and is used to group repeated runs of the same workload specification for statistical comparison.

The `config_hash()` is deliberately a workload hash rather than a full execution-context hash. Backend choice, AD mode, requested thread count, observed thread count, OpenMP/Rayon environment variables, software versions and hardware metadata are stored in separate database fields. The effective reproducibility key for a timed result is therefore the combination of `config_hash`, engine, AD mode, parallelism metadata and environment metadata, not `config_hash` alone.

All sources of randomness are explicitly seeded. For a given engine, machine, thread setting and software environment, the same effective configuration produces repeatable prices. Different engines use different RNG implementations (NumPy `default_rng`, JAX `PRNGKey`, C++ `mt19937_64`, Rust `SmallRng`, CUDA `curand`), so the same seed produces statistically equivalent but numerically distinct results across engines. In addition, the C++ and Rust engines use per-thread RNG streams, so changing the thread count changes the path-to-RNG-stream assignment and can change bit-level outputs even when the workload hash is unchanged. Cross-engine validation is therefore performed against analytical references and tolerances rather than against bit-equal outputs.

Extensibility. The framework is organised around two minimal abstractions: `WorkloadConfig` (what to simulate) and `MonteCarloEngine` (how to simulate it). Adding a workload requires only declaring a new `@workload`-decorated dataclass. Adding an engine requires only subclassing `MonteCarloEngine` and implementing `run()`, `supports()` and `supported_ad_modes()`. A central `WORKLOAD_REGISTRY` maps string identifiers to workload constructors and is consumed by both the REST API and the React frontend, so new workloads appear automatically in the UI without backend–frontend duplication.

Fairness. All engines are evaluated under a single benchmarking protocol implemented in `BenchmarkRunner`. Each trial consists of a configurable warmup phase whose results are discarded, followed by a fixed number of timed repetitions. Warmup is essential for JIT-compiled backends: a cold JAX call at $M = 10,000$ takes approximately 115 ms, compared to approximately 0.6 ms once XLA has cached the compiled binary. Excluding the cold compile cost from timing isolates the steady-state computational performance that is relevant in a production setting.

Traceability. Every row in the SQLite database carries the full serialised configuration (`config_json`), the analytical reference values, all timing statistics, and a complete snapshot of the execution environment (Python version, platform, CPU model and core count, RAM, NumPy

version, JAX version). Any historical run can be reconstructed exactly, even if the corresponding `WorkloadConfig` class is later modified.

3.2 High-Level Architecture

The framework follows a strict four-layer architecture, with two additional layers providing the REST API and the React frontend. Each layer depends only on the layers below it, so any component can be replaced or extended in isolation.




Figure 3.1: High-level system architecture and data flow.

3.3 Layer 1 — Workload Configuration

The workload layer uses a registry pattern implemented via a single Python decorator. The `@workload(name)` decorator simultaneously applies `@dataclass` to the class, sets a class-level `workload_type` attribute, and registers the class in a global `WORKLOAD_REGISTRY` dictionary. Adding a new workload therefore requires only writing a new class: no boilerplate, no changes to the runner, storage, API or frontend.

The base class `WorkloadConfig` provides four concrete inherited methods used by every workload:

- `validate()` — a no-op by default, overridden for cross-field constraints (e.g. enforcing $\sigma > 0$, $T > 0$, valid `option_type`).
- `to_dict()` — full serialisation via `dataclasses.asdict()` with `workload_type` injection.
- `from_dict()` — safe deserialisation, ignoring unknown keys so that schema migrations do not break older saved configurations.
- `config_hash()` — the 8-character SHA-256 fingerprint of the sorted serialised configuration.

There are no abstract methods on `WorkloadConfig`; every subclass inherits working implementations. The free function `config_from_dict()` is the single deserialisation entry point: it reads

the `workload_type` key, looks up the corresponding class in `WORKLOAD_REGISTRY`, and delegates to that class's `from_dict()`.

3.3.1 Workload 1 — European Option

`EuropeanOptionConfig` (`workload_type = "european"`) describes a plain vanilla European call or put under constant-volatility geometric Brownian motion (GBM). Its fields are summarised in Table 3.1.

Table 3.1: Fields of `EuropeanOptionConfig`.

Field	Default	Description
<code>S0</code>	100.0	initial spot price
<code>K</code>	100.0	strike price
<code>r</code>	0.05	continuously-compounded risk-free rate
<code>sigma</code>	0.20	annualised volatility
<code>T</code>	1.0	time to maturity (years)
<code>option_type</code>	"call"	"call" or "put"
<code>N</code>	1	number of time steps (single-step exact)
<code>M</code>	10 000	number of Monte Carlo paths
<code>seed</code>	42	RNG seed for reproducibility

The pricing model uses single-step exact GBM. The terminal stock price is

$$S_T = S_0 \exp\left(\left(r - \frac{1}{2}\sigma^2\right)T + \sigma\sqrt{T}Z\right), \quad Z \sim \mathcal{N}(0, 1),$$

and the option price is the discounted expected payoff. This is the reference workload for all engine and AD comparisons because Black–Scholes provides an exact analytical benchmark for both prices and Greeks.

Analytical validation is built directly into the framework. The `CPUMonteCarloEngine` module exposes Black–Scholes closed-form functions for the call/put price, Delta, Vega and Rho:

$$\Delta_{\text{call}} = \Phi(d_1), \quad \text{Vega} = S_0 \phi(d_1) \sqrt{T}, \quad \rho_{\text{call}} = K T e^{-rT} \Phi(d_2),$$

where $d_1 = (\ln(S_0/K) + (r + \frac{1}{2}\sigma^2)T)/(\sigma\sqrt{T})$ and $d_2 = d_1 - \sigma\sqrt{T}$. These values are computed and stored alongside every Monte Carlo result, enabling automatic absolute and relative error metrics.

3.3.2 Workload 2 — European Local-Volatility Option

`EuropeanLocalVolConfig` (`workload_type = "european_local_vol"`) extends the European workload to a multi-step local-volatility model. The configuration adds the parameters in Table 3.2.

Table 3.2: Additional fields of `EuropeanLocalVolConfig`.

Field	Default	Description
<code>M</code>	100 000	paths (larger than European for multi-step precision)
<code>N</code>	252	Euler steps (daily over one year)
<code>sigma_min</code>	0.01	volatility floor
<code>theta</code>	auto	polynomial parameters $[a_0, a_1, a_2, b_1]$

The default θ is set in `__post_init__` so that $\sigma_{\min} + \text{softplus}(a_0) = 0.20$, i.e. the model reduces to a 20% flat-volatility GBM. The closed-form solution is $a_0 = \log(\text{expm1}(0.20 - \sigma_{\min}))$. When $a_1 = a_2 = b_1 = 0$, the model collapses to constant-volatility GBM and the Monte Carlo price converges to the Black–Scholes price as M and N increase. This provides an automatic regression test embedded in the default configuration.

The dynamics are

$$dS = r S dt + \sigma(S, t; \theta) S dW,$$

discretised in log-space using a log-Euler scheme:

$$S_{n+1} = S_n \exp\left(\left(r - \frac{1}{2}\sigma_n^2\right)\Delta t + \sigma_n \sqrt{\Delta t} Z_n\right).$$

At each step the local volatility is computed from a parametric monomial surface:

$$x_n = \ln(S_n/S_0), \quad \text{raw}_n = a_0 + a_1 x_n + a_2 x_n^2 + b_1 t_n, \quad \sigma_n = \sigma_{\min} + \text{softplus}(\text{raw}_n).$$

The softplus activation is implemented in a numerically stable form $\max(\text{raw}, 0) + \log(1 + \exp(-|\text{raw}|))$, which avoids overflow for large positive inputs and catastrophic cancellation for large negative inputs. This formulation is reproduced *identically* in NumPy, JAX, C++ and Rust so that the engines are differentiated only by their execution model, not by numerical behaviour.

The AD targets for this workload are Delta ($\partial P/\partial S_0$) plus the full θ -gradient $\partial P/\partial a_0$, $\partial P/\partial a_1$, $\partial P/\partial a_2$, $\partial P/\partial b_1$, $\partial P/\partial \sigma_{\min}$ — six inputs in total, computed in a single reverse-mode backward pass.

3.4 Layer 2 — Engine Interface

The `MonteCarloEngine` abstract base class defines a minimal interface:

```
class MonteCarloEngine:
    def run(self, config, ad_mode) -> tuple[float, dict | None]
    def supports(self, workload_type: str) -> bool
    def supported_ad_modes(self) -> tuple[str, ...]
```

The `ad_mode` argument takes one of three string values: `"none"`, `"forward"` or `"reverse"`. When `ad_mode` is `"none"` or the engine does not implement AD for the given workload, the returned `greeks` dictionary is `None`. The runner inspects `supports()` and `supported_ad_modes()` before dispatching a configuration, so engines that cannot service a request are silently skipped rather than raising errors.

Engine 1 — NumPy / CPU. `CPUMonteCarloEngine` is the single-threaded NumPy reference implementation. It supports the implemented non-AD workloads and `ad_mode = "none"` only. For the European workload it uses `np.random.default_rng(seed)` to generate all M standard normals as a single array and evaluates the terminal price and payoff vectorially — no Python-level loop over paths. For the local-volatility workload it implements the N -step log-Euler scheme in Python, with each step operating vectorially on an array of M paths. The correctness of all other engines is verified against this implementation.

Engine 2 — JAX / XLA. `JAXMonteCarloEngine` is the only engine that exposes `ad_mode ∈ {none, forward, reverse}`. The European kernel is a *module-level* `@jax.jit` function decorated with `functools.partial(jax.jit, static_argnames=("option_type",))`. Defining the kernel at module scope ensures JAX traces and compiles it exactly once per `(shape, dtype, option_type)` signature; the random sample Z is passed as an argument rather than captured in a closure, so the compiled XLA computation is reused across every timed repetition. The local-volatility kernel uses `jax.lax.scan` to implement the N -step Euler loop, with `static_argnames=("option_type", "N")` so XLA emits specialised code for each call/put and step-count combination.

For the European workload, reverse-mode AD is implemented as a *single* `jax.grad(price_fn, argnums=(0,1,2))` call returning $(\partial P/\partial S_0, \partial P/\partial r, \partial P/\partial \sigma)$ in one backward pass. An earlier version invoked `jax.grad` three times with `argnums=0,1,2`, which was correct but tripled both compute time and peak memory — a result that materially distorts the AD overhead measurement. Forward-mode uses three `jax.jvp` calls probing the standard basis tangent vectors $(1, 0, 0)$, $(0, 1, 0)$ and $(0, 0, 1)$, exploiting the fact that forward mode is cheaper than reverse when the number of inputs is small relative to the number of outputs.

For the local-volatility workload, reverse-mode uses a single `jax.grad(..., argnums=(0,1,2,3,4,5))` call that differentiates through the full N -step `scan`; forward-mode uses six `jax.jvp` calls. JAX differentiates through `jax.lax.scan` efficiently via its reverse-mode accumulation, avoiding the explosion in graph size that would occur if the loop were unrolled.

Engine 3 — C++ / OpenMP. `CPPMonteCarloEngine` is a thin Python wrapper around a native implementation in `cpu_engine.cpp` / `cpu_engine.hpp`, exposed via `pybind11` bindings. Each OpenMP thread owns its own `std::mt19937_64` generator seeded as `seed + thread_id`; an

explicit `#pragma omp parallel` block with a reduction clause accumulates partial sums without locks. Results are fully deterministic for a given `(M, seed, thread_count)` triple; changing the thread count changes the per-path RNG assignment and produces statistically equivalent but numerically different results. The local-vol kernel reproduces exactly the same log-Euler discretisation and stable softplus formula as the NumPy and JAX engines.

Engine 4 — Rust / Rayon. `RustMonteCarloEngine` wraps a native implementation in `benchmarking/rust/src/lib.rs`, exposed via PyO3 bindings and built with `maturin`. The engine uses `SmallRng` (Xoshiro128++) from the `rand` crate, with each Rayon worker thread seeded as `seed XOR thread_index` to mirror the C++ per-thread offset pattern. Sampling uses `rand_distr::StandardNormal`. The per-path simulation is a pure scalar loop with no heap allocation; the entire M -path simulation is expressed as a Rayon `into_par_iter` chain with `map_init` for per-thread RNG initialisation and a `sum` reduction. The softplus is implemented inline as `z.max(0.0) + (-z.abs()).exp().ln_1p()`.

Engine 5 — CUDA / PyCUDA. `CUDAMonteCarloEngine` is included as European-only prototype GPU infrastructure. It compiles PyCUDA kernels at runtime, probes availability with a small startup run, and is excluded automatically when CUDA is unavailable. It supports no-AD European pricing and a pathwise European Delta kernel, but the final quantitative conclusions are based on the CPU engines and cloud CPU instances evaluated in Chapter 6.

3.5 Layer 3 — Benchmark Runner

`BenchmarkRunner` encapsulates the full benchmarking protocol. Given an engine instance and an optional display name, its `run()` method executes the following sequence:

1. Call `config.validate()`; abort early on failure.
2. Execute `num_warmup` warmup runs whose results are discarded; this excludes XLA compilation from timed measurements.
3. Execute `num_runs` timed runs, recording individual wall-clock times via `time.perf_counter()`.
4. If `ad_mode != "none"`, repeat steps 2–3 for a matched no-AD baseline (same engine, same config, identical warmup count) and compute the AD overhead ratio $r_{AD} = \bar{t}_{AD} / \bar{t}_{baseline}$.
5. Snapshot environment metadata: timestamp, Python version and implementation, platform, CPU model and architecture, logical CPU count, total RAM in GB, NumPy version, JAX version.
6. Aggregate timings into mean, standard deviation, min and max (in milliseconds), and compute `throughput_paths_per_sec = M / mean_runtime_s`.
7. Return a `BenchmarkResult` dataclass containing all of the above plus the price, the Greeks (if any), the analytical reference values, and the absolute and relative errors.

Memory peak is measured via `psutil`: the process resident-set size is snapshotted before and after the timed run set, and the peak is reported as $\max(rss_{during}) - rss_{before}$. If `psutil` is unavailable the field is left blank, degrading gracefully.

3.6 Layer 4 — Storage

`BenchmarkDB` is a thread-safe SQLite wrapper. It uses a per-call connection with `check_same_thread=False` and an explicit `threading.Lock` that serialises writes. The database is opened in Write-Ahead Logging (WAL) mode, which allows concurrent reads during writes. This matters because the Flask API serves read requests from the main thread while the background worker writes results. The database file lives at `results/benchmarks.db` and is created on first use.

Schema. The schema consists of a single `runs` table whose columns are defined declaratively by the `_FULL_SCHEMA_COLUMNS` list. Columns are grouped as follows:

- **Identity:** `id` (UUID, primary key), `experiment_id`, `experiment_type`, `config_hash`.
- **Status lifecycle:** `status` (pending/running/completed/failed), `error_message`, `created_at`, `started_at`, `completed_at` (all UTC ISO-8601).
- **Workload definition:** `workload_type`, `M`, `N`, `seed`, `config_json`.
- **Engine / implementation:** `engine`, `language`, `backend`.
- **Parallelism metadata:** `num_threads`, `vectorization_flag`, `batch_size`.
- **Runtime performance:** `mean_runtime_ms`, `std_runtime_ms`, `min_runtime_ms`, `max_runtime_ms`, `throughput_paths_per_sec`.
- **AD metadata:** `ad_mode`, `baseline_mean_ms`, `ad_overhead_ratio`.
- **Numerical results:** `result_value`, `greek_delta`, `greek_vega`, `greek_rho`.
- **Analytical reference:** `analytical_price`, `analytical_delta`, `analytical_vega`, `analytical_rho`.
- **Error metrics:** `abs_price_error`, `rel_price_error` and matching columns for Delta, Vega and Rho.
- **Resources:** `memory_peak_mb`.
- **Environment:** `cpu_model`, `cpu_architecture`, `cpu_count`, `memory_gb`, `platform`, `python_version`, `numpy_version`, `jax_version`, `blas_backend`.
- **Cloud and profiling:** `cloud_provider`, `instance_type`, `cloud_region`, `cloud_zone`, `hourly_price_usd`, `cost_per_run`, profiler-selection fields, and SHA promotion metadata where available.

Auto-migration. On startup, `_init_schema()` compares the existing table columns against the authoritative list using `PRAGMA table_info` and issues an `ALTER TABLE ADD COLUMN` for any missing column. Existing databases therefore continue to work unmodified as the schema evolves — a property that has proved valuable while iterating on the AD and cloud columns during development.

Write paths. Experiment scripts use `store_run_full()`, which inserts a fully-populated completed row in a single transaction, bypassing the pending/running state machine. The API uses the `create_run()` → `mark_running()` → `mark_completed()` lifecycle to support asynchronous execution.

3.7 Layer 5 — REST API

The Flask server (`benchmarking/api/server.py`, port 5050) is intentionally small. Flask-CORS is enabled for cross-origin requests from the frontend. A single daemon background thread polls `BenchmarkDB` for pending runs and executes them sequentially via `BenchmarkRunner`, transitioning the row through the `running` and `completed` (or `failed`) states. The endpoints are summarised in Table 3.3.

Table 3.3: REST API endpoints.

Endpoint	Method	Purpose
<code>/api/runs</code>	POST	submit a new benchmark run
<code>/api/runs</code>	GET	list runs (filter by workload/engine/status)
<code>/api/runs/:id</code>	GET	poll a single run by UUID
<code>/api/engines</code>	GET	list available engines, workloads and AD modes
<code>/api/workloads</code>	GET	list registered workload types
<code>/api/summary</code>	GET	aggregate counts, fastest run, breakdowns
<code>/api/compare_matrix</code>	GET	matrix aggregated by (engine, ad_mode, config)

Engine availability is determined at startup. The C++ and CUDA engines may not be present in every environment; they are excluded from `/api/engines` without erroring if their initialisation probe fails.

3.8 Layer 6 — Frontend

The frontend is a React + TypeScript single-page application built with Vite and styled with Material UI, listening on port 5173. It communicates with the Flask backend through a typed HTTP client in `src/api/client.ts`. The application has five pages:

- **SimulatePage** — submit a new run and poll for completion. The `SimulationForm` component is generated dynamically from the workload `SCHEMA` returned by the API, so new workloads appear automatically without frontend changes.
- **ComparePage** — engine comparison matrix backed by `/api/compare_matrix`.
- **HistoryPage** — filterable table of every run.
- **SummaryPage** — aggregate dashboard.
- **AnalysisPage** — AD overhead and Greek-accuracy analysis.

3.9 Cloud Cost-Performance Design

Cloud resource selection is treated as a practical systems question rather than a purely local benchmark. The final design therefore extends each benchmark row with cloud metadata whenever the run is executed on a cloud instance. The relevant fields record the provider, instance type, region, zone, hourly price assumption, git commit, and derived cost per run. This keeps cost-performance analysis tied to the exact environment in which a run was measured.

The cost model is intentionally simple:

$$\text{cost per run} = \frac{\text{runtime seconds}}{3600} \times \text{hourly instance price}.$$

This does not attempt to model billing granularity, storage, data transfer, or queueing costs. Its purpose is to compare benchmark configurations under a common pricing assumption. This is sufficient for identifying whether a faster instance is also cheaper per run for the measured workload, while avoiding unsupported claims about long-term cloud bills or all future machine prices.

Cloud fields are stored alongside the same workload, engine, timing, numerical, and environment metadata as local runs. As a result, cloud and local experiments share the same analysis path: report artefacts can be regenerated from the database without manually copying console output into tables.

3.10 Budget-Aware Profiler Design

Full-grid benchmarking gives the cleanest reference point, but it scales poorly with the number of workloads, engines, AD modes, path counts, and instance types. The framework therefore includes profiler logic that uses cheaper probe runs to select a subset of configurations for full-scale benchmarking. The important design principle is that the profiler is not trusted by construction: it is evaluated against a full-grid baseline using explicit selection-quality metrics.

3.10.1 Probe-Based Baseline

The first profiler uses low-budget probes to rank all candidate configurations and then selects a conservative subset for full benchmarking. It includes diversity safeguards so that selection is not determined only by a single global top- k ranking. This makes it a useful baseline: it can save full runs, but it still evaluates a relatively large selected subset to protect against probe noise and workload-specific crossover behaviour.

3.10.2 Successive Halving Profiler

The successive-halving (SHA) profiler evaluates all candidate configurations at the cheapest probe budget, promotes a smaller active set to a larger probe budget, and repeats this process before choosing the final full-scale runs. In the final cloud experiment, the budget is the number of Monte Carlo paths, so the profiler asks whether small-path measurements are sufficiently predictive of full-path performance.

The design deliberately records intermediate SHA state rather than only the final selected set. For each promotion stage, the database or generated artefacts capture the active configuration count and probe budget. These records allow the report to show how many candidates are discarded at each stage and to measure the probe path budget consumed by the profiler.

Selection quality is measured using:

- **full runs saved:** the reduction in expensive full-scale benchmark evaluations relative to the full grid;
- **Pareto recovery:** the fraction of full-grid non-dominated configurations retained by the profiler;
- **runtime regret:** the percentage loss relative to the fastest full-grid configuration;
- **cost regret:** the percentage loss relative to the cheapest full-grid configuration; and
- **rank correlation:** Spearman agreement between probe rankings and full-scale rankings.

These metrics make the profiler result falsifiable. A profiler that saves runs but misses the Pareto frontier or introduces high regret would not be considered successful. Conversely, a profiler that preserves the full-grid decisions while reducing full runs contributes directly to the cloud resource-selection goal of the project.

The final revision adds guarded SHA for the local-volatility deployment task. The guarded version keeps the efficient SHA structure, but separates price-only and AD-required candidate pools and applies explicit near-tie, engine-diversity, instance-family, scaling-improvement, and correctness/AD guards before final elimination. This design responds directly to the observed risk that a plain small-budget SHA ranking can discard configurations whose larger-budget scaling or cost-performance later improves.

3.11 Design Decisions and Trade-offs

Decorator-based workload registration. An earlier design made `WorkloadConfig` an abstract base class with abstract properties (`workload_type`, `M`, `seed`) and abstract methods (`validate()`, `to_dict()`, `from_dict()`). That design forced every workload to re-implement the same boilerplate and produced a class of bugs in which abstract method implementations drifted out of sync. The refactored design replaces all of this with a single `@workload(name)` decorator that wires up dataclass conversion, sets `workload_type`, and registers the class. New workloads need only declare their fields and optionally override `validate()`.

Module-level JAX kernels. Placing JAX kernels at module level (rather than as instance methods or nested functions) is critical for compilation caching. JAX traces and compiles based on function identity and on argument shapes and dtypes. If the kernel were a nested function inside `run()`, each call to `run()` would create a fresh function object and defeat the JIT cache, regenerating the XLA binary on every timed repetition.

Single-pass reverse-mode AD. Using a single `jax.grad(price_fn, argnums=(0,1,2))` call rather than three separate `jax.grad` calls is functionally equivalent but materially changes the measured AD overhead. The single-pass form shares the backward graph across all output partial derivatives and is consistent with how AD is deployed in production quant systems; this is the form benchmarked in this project.

jax.lax.scan for the local-volatility loop. The N -step Euler loop must be expressed as `jax.lax.scan` rather than as a Python `for` loop. A Python loop would cause JAX to unroll the entire computation during tracing, producing an XLA graph with $N = 252$ copies of the step body — both compile time and memory consumption grow with N . `jax.lax.scan` compiles to a single loop in the XLA graph and lets the compiler treat it as a unit.

SQLite over PostgreSQL. SQLite with WAL is sufficient for the expected single-machine workload (up to a few hundred runs per day). It requires no external service, ships as part of Python, and is straightforward to back up by copying a single file. WAL allows concurrent reads during writes, which is the only concurrency requirement imposed by the architecture. If the framework were scaled beyond a single-machine execution model, PostgreSQL would be a natural replacement.

Per-thread RNG streams. The C++ and Rust engines seed each worker thread independently. This sacrifices bitwise reproducibility across thread counts in exchange for lock-free parallel sampling. The framework records `num_threads` in the database so that runs at different thread counts are treated as statistically distinct execution settings rather than as bitwise repeats of the same run.

3.12 Summary

The framework consists of six clearly separated layers: workload configuration, engine, runner, storage, API and frontend. The four implementation layers are strictly decoupled through two minimal abstractions (`WorkloadConfig` and `MonteCarloEngine`) and a registry-driven workload model. Every benchmark run is fully described by a serialisable configuration, fingerprinted by a deterministic hash, and persisted alongside its analytical reference values and execution environment, ensuring complete reproducibility and traceability. The next chapter describes how each layer is implemented in practice.

Chapter 4

Implementation

4.1 Overview

This chapter describes the concrete implementation of the framework laid out in Chapter 3. The codebase is hosted in the project GitHub repository and is organised as a single Python package with optional native subpackages. The top-level directory structure is:

```
benchmarking/  
  workloads/      # WorkloadConfig dataclasses + registry  
  engines/        # Python wrappers for each engine  
  cpp/            # native C++ implementation (pybind11)  
  rust/           # native Rust implementation (PyO3)  
  runner.py       # BenchmarkRunner  
  db.py           # BenchmarkDB (SQLite)  
  api/            # Flask REST API  
frontend/         # React + TypeScript + Vite  
experiments/      # CLI experiment drivers  
results/          # SQLite database, generated tables  
tests/            # pytest suite
```

The Python package is pure-Python apart from optional native extensions. The C++ extension is built with `pipinstall-ebenchmarking/cpp/`, using a `setup.py` file that invokes `pybind11`. The Rust extension is built with `maturindevelop--release` inside `benchmarking/rust/`. If either extension is not built, the corresponding engine is silently excluded from the engine registry at runtime, so the framework remains fully functional on systems without a C++ or Rust toolchain.

4.2 Core Modules

4.2.1 Workload Registry

The workload registry is implemented in `benchmarking/workloads/base.py`. The decorator `@workload(name)` performs three actions in a single pass: it applies `dataclass` to the decorated class, sets a class-level attribute `workload_type = name`, and registers the resulting class in the module-level dictionary `WORKLOAD_REGISTRY`. The base class `WorkloadConfig` provides concrete implementations of `validate()`, `to_dict()`, `from_dict()` and `config_hash()` so that no subclass is required to override anything except its field list and (optionally) its validation logic.

The free function `config_from_dict(d)` reads `d["workload_type"]`, looks up the corresponding class in `WORKLOAD_REGISTRY`, and delegates to that class's `from_dict()`. Unknown keys are silently dropped during deserialisation, so configurations stored by older versions of the framework continue to load even after new fields are added. `config_hash()` serialises the configuration with `json.dumps(..., sort_keys=True)` and returns the first eight characters of the SHA-256 digest as a stable, machine-independent workload fingerprint. It does not include execution-context fields such as engine, AD mode, requested thread count or OpenMP/Rayon environment variables; those are recorded separately in the result database and must be considered alongside the hash when reproducing a specific timed run.

4.2.2 Monte Carlo Engine Interface

All engines are subclasses of the abstract base class `MonteCarloEngine`, defined in `benchmarking/engines/base.py`:

```
class MonteCarloEngine:
    name: str
    language: str
    backend: str

    def supports(self, workload_type: str) -> bool: ...
    def supported_ad_modes(self) -> tuple[str, ...]: ...
    def run(self, config: WorkloadConfig,
            ad_mode: str = "none") -> tuple[float, dict | None]: ...
```

The class attributes `name`, `language` and `backend` are the values written to the corresponding columns of `benchmarks.db`, enabling later queries such as “all Rust runs” or “all GPU runs”. The `run()` method is the only computational entry point: it accepts a fully-validated `WorkloadConfig` and an AD mode string, and returns (price, greeks). The `greeks` component is either `None` (for `ad_mode = "none"`) or a dictionary mapping Greek names (`"delta"`, `"vega"`, `"rho"`, ...) to floats.

4.2.3 NumPy Engine

`CPUMonteCarloEngine` (`benchmarking/engines/mc_cpu.py`) is the single-threaded NumPy reference implementation. For the European workload it draws all M standard normals in one call to `np.random.default_rng(seed).standard_normal(M)`, evaluates the closed form for S_T vectorially and computes the discounted mean payoff. For the local-volatility workload it runs the N -step log-Euler loop in Python, with each iteration operating on a NumPy array of M paths and using the stable softplus

```
def softplus(z):
    return np.maximum(z, 0.0) + np.log1p(np.exp(-np.abs(z)))
```

The same module exposes Black–Scholes closed-form helpers for price, Delta, Vega and Rho, which the runner calls to populate the analytical reference columns of every database row.

4.2.4 JAX Engine

`JAXMonteCarloEngine` (`benchmarking/engines/mc_jax.py`) is the only engine that implements AD. The European kernel is declared at module level to maximise compilation cache hits:

```
@functools.partial(jax.jit, static_argnames=("option_type",))
def _european_price(S0, r, sigma, K, T, Z, option_type):
    sqrtT = jnp.sqrt(T)
    ST = S0 * jnp.exp((r - 0.5 * sigma**2) * T + sigma * sqrtT * Z)
    payoff = jnp.where(option_type == "call",
                       jnp.maximum(ST - K, 0.0),
                       jnp.maximum(K - ST, 0.0))
    return jnp.exp(-r * T) * jnp.mean(payoff)
```

Reverse-mode Greeks for the European workload are obtained from a single backward pass:

```
grad_fn = jax.jit(jax.grad(_european_price, argnums=(0, 1, 2)),
                  static_argnames=("option_type",))
delta, rho, vega = grad_fn(S0, r, sigma, K, T, Z, option_type)
```

Forward-mode uses three `jax.jvp` calls, one per input direction, each probing the corresponding basis tangent vector. The local-volatility kernel expresses the N -step loop as `jax.lax.scan`:


```

def _step(carry, z):
    S, t = carry
    x = jnp.log(S / S0)
    raw = a0 + a1 * x + a2 * x * x + b1 * t
    sigma = sigma_min + softplus(raw)
    S_next = S * jnp.exp((r - 0.5 * sigma**2) * dt
                        + sigma * jnp.sqrt(dt) * z)
    return (S_next, t + dt), None

(S_T, _), _ = jax.lax.scan(_step, (S0_arr, 0.0), Z_matrix)

```

`jax.lax.scan` produces a single loop in the XLA graph; tracing time and peak memory remain constant in N , in contrast with a Python `for` loop which would unroll the graph to size $\Theta(N)$ and substantially increase compile time and memory for $N = 252$.

The AD overhead is computed by the runner as the ratio of the mean timed runtime of the AD run to that of an identically-configured no-AD baseline run. Crucially, both legs use the same warmup count, so XLA cache state is equivalent and the ratio reflects only the differentiation cost.

4.2.5 C++ Engine

The native implementation lives in `benchmarking/cpp/cpu_engine.cpp / cpu_engine.hpp` and is exposed to Python via `pybind11` bindings in `benchmarking/cpp/bindings/pybind_module.cpp`. The Python wrapper `benchmarking/engines/mc_cpp.py` loads the compiled module lazily and falls back to a not-available stub if the import fails.

Parallelisation is via OpenMP. Each thread owns its own `std::mt19937_64` generator, seeded as `seed + thread_id` to guarantee statistical independence across threads while remaining reproducible for any fixed thread count. The European inner loop is:

```

double sum = 0.0;
#pragma omp parallel reduction(+:sum)
{
    int tid = omp_get_thread_num();
    std::mt19937_64 rng(seed + tid);
    std::normal_distribution<double> N01(0.0, 1.0);
    #pragma omp for
    for (int i = 0; i < M; ++i) {
        double Z = N01(rng);
        double ST = S0 * std::exp((r - 0.5 * sigma * sigma) * T
                                + sigma * std::sqrt(T) * Z);
        sum += std::max(option_type == 0 ? ST - K : K - ST, 0.0);
    }
}
double price = std::exp(-r * T) * sum / M;

```

The local-volatility kernel reuses the same RNG strategy and implements the log-Euler loop with the same numerically stable `softplus` used by the Python and Rust formulations:

$$\max(\text{raw}, 0) + \log(1 + \exp(-|\text{raw}|)).$$

4.2.6 Rust Engine

The native implementation lives in `benchmarking/rust/src/lib.rs` and is exposed to Python via `PyO3` bindings, built with `maturin develop --release`. The Python wrapper `benchmarking/engines/mc_rust.py` mirrors the C++ wrapper: lazy import, graceful fallback, identical engine interface.

The Rust implementation uses `SmallRng` (Xoshiro128++) from the `rand` crate for high-throughput sampling, with each Rayon worker thread seeded as `seed XOR thread_index`. The simulation is expressed as a data-parallel iterator over the path range:

```

let price: f64 = (0..M).into_par_iter()
    .map_init(
        || SmallRng::seed_from_u64(seed ^ rayon::current_thread_index()
                                     .unwrap_or(0) as u64),

        |rng, _| {
            let z: f64 = StandardNormal.sample(rng);
            let s_t = s0 * ((r - 0.5 * sigma * sigma) * t
                           + sigma * t.sqrt() * z).exp();
            (s_t - k).max(0.0)
        })
    .sum:<f64>() * (-r * t).exp() / (M as f64);

```

The softplus is implemented inline as `z.max(0.0) + (-z.abs()).exp().ln_1p()`, matching the Python, C++ and JAX formulations bit-for-bit.

4.2.7 CUDA Engine

`CUDAMonteCarloEngine` (`benchmarking/engines/mc_cuda.py`) uses PyCUDA to compile and launch CUDA kernels at runtime. At import time, the engine runs a minimal 128-path trial; if CUDA initialisation, kernel compilation, or the trial run fails, the engine sets its availability flag to `False` and is excluded from `/api/engines`.

The European kernel uses one thread per Monte Carlo path:

```

__global__ void price_kernel(double S0, double r, double sigma,
                             double K, double T, double sqrtT,
                             unsigned long long seed,
                             int option_type, int M,
                             double *payoffs) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx >= M) return;
    curandState state;
    curand_init(seed, idx, 0, &state);
    double Z = curand_normal_double(&state);
    double ST = S0 * exp((r - 0.5 * sigma * sigma) * T + sigma * sqrtT * Z);
    double payoff = option_type == 0
        ? fmax(ST - K, 0.0) : fmax(K - ST, 0.0);
    payoffs[idx] = exp(-r * T) * payoff;
}

```

Using the thread index as the `curand_init` sequence guarantees statistical independence regardless of block or grid layout. The host averages the resulting payoff array with `np.mean`.

Pathwise forward-mode Delta is implemented in a second kernel. Using the identity $\partial S_T / \partial S_0 = S_T / S_0$ for GBM, the pathwise Delta integrand is

$$\mathbf{1}_{\{S_T > K\}} \cdot \frac{S_T}{S_0} \quad (\text{call}), \quad -\mathbf{1}_{\{S_T < K\}} \cdot \frac{S_T}{S_0} \quad (\text{put}),$$

and the engine returns $\Delta = e^{-rT} \mathbb{E}[d(\text{payoff})/dS_0]$. Note that this is a *pathwise* estimator, not a finite-difference approximation; it produces unbiased Greeks with the same variance characteristics as a single-pass forward-mode AD.

4.2.8 Benchmark Runner

`BenchmarkRunner` (`benchmarking/runner.py`) is a small class. Construction takes an engine instance, an optional display name and the warmup and repetition counts. The `run()` method implements the protocol of Section 3.5. Timed repetitions use `time.perf_counter()` which provides nanosecond resolution on Linux and macOS. The runner separately computes the analytical reference values (by calling the Black–Scholes helpers in `mc_cpu.py`) and the absolute and relative errors for the price and each Greek, so that the database row produced is self-contained.

When `ad_mode != "none"` the runner first performs the AD timed run, then a matched no-AD baseline run using the same engine, same configuration, same warmup count and same repetition count. This ensures that the AD overhead ratio reflects only the cost of differentiation, not the cost of XLA recompilation or any other one-time setup.

4.2.9 Metrics Collection

The runner collects four classes of metrics:

- **Timing:** mean, standard deviation, min, max of the timed repetitions, in milliseconds, plus the derived `throughput_paths_per_sec = M / mean_runtime_s`.
- **AD overhead:** the matched baseline mean runtime (`baseline_mean_ms`) and the overhead ratio $\bar{t}_{AD}/\bar{t}_{baseline}$ (`ad_overhead_ratio`).
- **Memory:** process RSS sampled before and after the timed repetitions via `psutil`; the recorded value is the difference.
- **Numerical accuracy:** absolute and relative error for the price and for each Greek against the Black-Scholes reference, when applicable.

Environment metadata is captured once per `run()` call: the wall-clock timestamp, the Python version and implementation, the platform string, the CPU model and architecture, the logical CPU count, total RAM in GB, the NumPy version and the JAX version.

4.2.10 Storage Layer

BenchmarkDB (`benchmarking/db.py`) is a thin SQLite wrapper. A new connection is opened for each call with `check_same_thread=False`, and an explicit `threading.Lock` serialises write operations. The database is created in WAL mode the first time it is opened. The schema is described in detail in Section 3.6; it is defined declaratively in the `_FULL_SCHEMA_COLUMNS` constant, and any missing columns are added automatically at startup via `PRAGMA table_info` followed by `ALTER TABLE`.

Two write paths coexist. Experiment scripts use `store_run_full()`, which writes a fully-populated `completed` row in a single transaction. The REST API uses the three-step lifecycle `create_run()` → `mark_running()` → `mark_completed()` (or `mark_failed()`) to support asynchronous execution from a background worker thread.

4.2.11 REST API

The Flask application (`benchmarking/api/server.py`) registers the seven endpoints listed in Table 3.3 and starts a single daemon background thread on startup. The worker polls the `runs` table for `pending` rows in FIFO order, dispatches each through the appropriate engine via `BenchmarkRunner.run()`, and updates the row with the result. Engine availability is probed once at startup: C++ and CUDA engines that fail to load or fail their probe runs are removed from the registry.

The API returns JSON in all cases and uses standard HTTP status codes. Flask-CORS is configured to allow requests from the Vite development server (port 5173) so that the frontend can be developed locally without a reverse proxy.

4.2.12 Frontend / Dashboard

The frontend is a React + TypeScript single-page application built with Vite and styled with Material UI. The directory layout follows the conventional `src/pages/`, `src/components/` and `src/api/` structure. A typed HTTP client in `src/api/client.ts` wraps every API endpoint with a TypeScript function whose return type mirrors the corresponding Python dataclass.

Pages:

- **SimulatePage** — submits a new run via `POST /api/runs` and polls `GET /api/runs/:id` until the status becomes `completed` or `failed`.
- **ComparePage** — consumes `/api/compare_matrix` to render a multi-engine, multi-AD-mode comparison grid.

- `HistoryPage` — filterable table backed by GET `/api/runs`.
- `SummaryPage` — consumes `/api/summary`.
- `AnalysisPage` — focused view of AD overhead and Greek accuracy.

Components are factored to be workload-agnostic: `SimulationForm` is generated dynamically from the workload `SCHEMA` returned by the API, so adding a new workload class (or even a new field) requires *no* frontend changes.

4.3 Supported Workloads

Three workloads are implemented, with validation strength depending on the available reference:

- `european` (`EuropeanOptionConfig`): single-step exact GBM. Supported by all five engines. Used for AD overhead, Greek accuracy, language comparison and throughput scaling experiments.
- `european_local_vol` (`EuropeanLocalVolConfig`): N -step log-Euler simulation of GBM with a parametric local-volatility surface. Evaluated through NumPy, JAX, C++ and Rust in the completed experiments. Used to stress-test AD on a deeper computation graph and to expose the cost of multi-step over single-step simulation. Its prices are checked by cross-engine agreement and convergence rather than a closed-form reference, with the Rust discrepancy treated as a validation caveat.
- `asian` (`AsianOptionConfig`): arithmetic-average path-dependent option workload. Supported by the Python and JAX engines in the final cloud-profiler experiments. Used to check that backend rankings and AD overheads are not inferred only from terminal-payoff European workloads.

4.4 AD Integration

Automatic differentiation is implemented exclusively in the JAX and CUDA engines. The JAX engine supports both forward and reverse mode for the European, local-volatility, and Asian workloads used in the final results; the CUDA engine implements only forward-mode pathwise Delta on the European workload and is not part of the final cloud comparison.

Table 4.1 makes the implemented workload and AD support explicit. Each cell lists the AD modes available for that workload and engine; a dash denotes an unsupported combination that the runner skips before dispatch.

Table 4.1: Implemented workload and AD support by engine.

Workload	NumPy	JAX	C++	Rust	CUDA
European GBM	none	none, forward, reverse	none	none	none, forward
Local volatility	none	none, forward, reverse	none	none	–
Asian arithmetic-average	none	none, forward, reverse	–	–	–

For the European workload the AD output is a three-tuple

$$(\Delta, \rho, \text{Vega}) = \left(\frac{\partial P}{\partial S_0}, \frac{\partial P}{\partial r}, \frac{\partial P}{\partial \sigma} \right).$$

Forward mode is implemented as three `jax.jvp` calls (one per input direction), and reverse mode as a single `jax.grad` call with `argnums=(0,1,2)`.

For the local-volatility workload the AD output is a six-tuple

$$\left(\frac{\partial P}{\partial S_0}, \frac{\partial P}{\partial a_0}, \frac{\partial P}{\partial a_1}, \frac{\partial P}{\partial a_2}, \frac{\partial P}{\partial b_1}, \frac{\partial P}{\partial \sigma_{\min}} \right).$$

Forward mode therefore requires six `jax.jvp` calls, while reverse mode again computes the full gradient in a single backward pass.

For the Asian workload, JAX differentiates the arithmetic-average payoff with respect to the same scalar market parameters used by the European GBM configuration. Both AD modes feed into the runner’s overhead-ratio computation. Because the gradient and the baseline price share the same forward computation, the overhead ratio is a clean measure of the marginal cost of differentiation.

4.5 GPU Support

CUDA support is implemented as prototype infrastructure for the European workload via PyCUDA. The engine launches one path per CUDA thread, uses cuRAND for normal draws, writes path payoffs to device memory, and reduces on the host. It is probed at server startup with a 128-path trial run and excluded if CUDA or PyCUDA is unavailable. JAX can also execute on GPU when installed with CUDA wheels. The final quantitative conclusions in Chapter 6, however, are based on the CPU backends and cloud CPU instances evaluated there; GPU evaluation is left as future work.

4.6 Challenges

Several issues arose during implementation and shaped the final design.

XLA caching across runs. A naive initial JAX implementation defined the kernel inside the engine’s `run()` method. Each call created a fresh function object, defeating the JIT cache: every timed repetition paid the full ≈ 115 ms compile cost. Moving the kernel to module level reduced warm execution time to ≈ 0.6 ms and made the JAX engine genuinely competitive on small M .

Triple AD passes. An earlier reverse-mode implementation invoked `jax.grad` three times with `argnums=0,1,2`. The result was numerically correct but tripled both the measured AD wall time and the AD overhead ratio. Consolidating into a single `jax.grad(..., argnums=(0,1,2))` call eliminated this artefact; this single change halved the reverse-mode overhead in early measurements and brought the framework into line with how reverse-mode AD is deployed in practice.

for-loop unrolling in JAX. The first local-volatility implementation used a Python `for` loop over the $N = 252$ Euler steps. Compilation time grew linearly in N and peak memory reached several gigabytes during tracing. Rewriting the loop as `jax.lax.scan` produced a single fused loop in the XLA graph; compile time and memory became effectively independent of N .

Numerically stable softplus. A naive softplus implementation $\log(1 + e^z)$ overflows for $z \gtrsim 700$ in double precision and loses precision near $z = 0$. The stable form $\max(z, 0) + \log(1 + e^{-|z|})$ is used identically in NumPy, JAX, C++ and Rust, so any cross-engine numerical discrepancy can only be attributed to the RNG or the floating-point order of operations, not to the activation function itself.

Per-thread RNG reproducibility. Both the C++ and Rust engines use per-thread RNG streams seeded from the global seed plus the thread index. Changing the thread count therefore changes the RNG assignment and produces statistically equivalent but numerically different results. This is recorded by storing `num_threads` on every run, so that runs at different thread counts are treated as distinct execution settings even when the workload-level `config_hash` is unchanged.

Optional native components. Both the C++ and Rust engines require non-trivial native toolchains (a C++17 compiler with OpenMP, or the Rust toolchain plus maturin) that are not present in every environment. The engines are loaded lazily and degrade gracefully: a missing build is reported only when the user explicitly requests that engine via the API, and the corresponding rows simply do not appear in the engine list.

Schema evolution. During development, several new columns were added to the database (notably the AD-overhead, error-metric and cloud columns). Rather than write migration scripts, the database opens with a self-migrating `_init_schema()` that compares the existing columns against the authoritative list and `ALTER TABLE`s any missing columns into existence. This has allowed every historical database file to remain readable across schema changes.

4.7 Cloud Metadata and Cost Support

Cloud support is implemented as metadata attached to ordinary benchmark rows rather than as a separate execution path. Experiment drivers can record the cloud provider, instance type, region, zone, hourly price assumption, git commit, and derived cost per run. This preserves the same framework invariants as local execution: every cloud result still has a workload configuration, an engine, an AD mode, timing statistics, numerical outputs, and environment metadata.

The final cloud experiment uses Google Cloud instance metadata recorded in the result files, including `n2-standard-4` in `europa-west1-b`. Cost per run is derived from measured runtime and the stored hourly price. The implementation deliberately avoids hard-coding a claim that one instance is universally cheapest: prices and availability can change, so the report treats cost as experiment metadata associated with a particular run.

4.8 Profiler Implementation

The final profiler-versus-grid experiment is driven by `experiments/run_profiler_vs_grid.py`. The script constructs the candidate grid, filters unsupported workload/engine/AD combinations, executes probe runs, applies the old profiler and SHA selection policies, and records which configurations would be promoted to full benchmarking. The report artefact generator `Final_Year_Project_Report/evaluation/generate_report_assets.py` then converts the stored results into the tables and figures used in Chapter 6.

In the final `sha_cloud_profiler_v4` experiment, the full-grid baseline contains 16 candidate configurations evaluated at three full path counts, for 48 full benchmark cells. The old probe-based profiler selects 30 full cells, while SHA selects 24. The SHA implementation records intermediate promotion rounds, including the active configuration count at each probe path budget. This metadata is used to generate the SHA progression plot and to compute probe-budget savings.

The profiler implementation also computes decision-quality metrics against the full grid: Pareto-frontier recovery, runtime regret, cost regret, and Spearman rank correlation between probe and full-scale rankings. This is important because the profiler is an optimisation of the benchmarking process, not a replacement for evaluation. Its success is judged by whether it reduces expensive full runs while preserving the decisions that the full grid would have supported.

The guarded local-volatility profiler is implemented in a dedicated `guarded_sha` package and executed by `run_guarded_sha_local_vol.py`. This keeps the new policy reusable rather than embedding it in the older monolithic profiler script. The implementation defines typed candidate, objective and guard objects, filters unsupported engine-AD combinations using engine capability metadata, records promotion/elimination reasons, and writes JSON, CSV and Markdown artefacts. A GCP wrapper script runs the same experiment across the selected instance types and then performs a combined cross-instance analysis locally.

4.9 Testing

The framework ships with 55 tests across two pytest files in `tests/`.

`test_system.py` covers:

- Six parametric tests for configuration validation (invalid field values, rejected `option_type`).
- Configuration serialisation: round-trip `to_dict` / `from_dict`, exclusion of `SCHEMA` from the serialised form, `config_from_dict` dispatch, and the expected `ValueError` on an unknown `workload_type`.
- Engine correctness: NumPy and JAX pricing accuracy against Black-Scholes within tolerance.
- Runner statistics: structure of the timing record and reproducibility under fixed seed.

- Database round-trips: `create_run`, `mark_running`, `mark_completed`, `get_run`, `get_all_runs`, `summary`.
- API endpoint contracts: status codes and response structure for all seven endpoints.
- An end-to-end flow from configuration through engine and runner to database insert and API read-back.
`test_ad_framework.py` covers:
 - JAX engine executes without error for each of `ad_mode` $\in$ `{none, forward, reverse}` on the European workload.
 - The standalone differentiated price function exposes the expected JAX function signature.
 - Analytical Greek computations agree with their closed-form definitions.
 - Computed gradients fall within the expected analytical bounds.

4.10 Dependencies

Core Python dependencies: `numpy>=1.26`, `scipy>=1.11`, `jax[cpu]>=0.4.28`, `flask>=3.0`, `flask-cors>=4.0`, `pytest>=8.0`, `pybind11>=2.11` (for the C++ build). `psutil` is optional and enables RSS-based memory profiling when present. The Rust engine requires the Rust toolchain (`rustup`), `maturin>=1.0`, `PyO3`, and the `rand`, `rand_distr` and `rayon` crates. The CUDA engine requires the NVIDIA CUDA toolkit and `PyCUDA`. The frontend depends on `react`, `react-dom`, `@mui/material` and a Vite-based toolchain.

4.11 Summary

The implementation realises the design of Chapter 3 in roughly 4,000 lines of Python plus native C++ and Rust extensions and a TypeScript frontend. The engine registry includes NumPy, JAX, C++, Rust, and CUDA paths, while the completed experiments focus on the supported NumPy, JAX, C++ and Rust combinations. Three workloads (European GBM, European local volatility, and Asian arithmetic-average options) are represented in the framework. The no-AD mode is available across the relevant engine set, while forward and reverse AD are evaluated through JAX for the completed results, with a narrower European-only CUDA AD path implemented but not used in the final cloud comparison. The system also includes full SQLite persistence, a REST API, a React dashboard and a 55-test pytest suite. The next chapter describes the experimental protocol applied to this implementation.

Chapter 5

Experimental Setup

5.1 Overview

This chapter defines the experimental methodology used to evaluate the Monte Carlo benchmarking framework. The aim is to ensure that all measurements reported in Chapter 6 are produced from controlled, reproducible, and scientifically interpretable experiments.

The evaluation combines two forms of evidence. First, a detailed local CPU campaign is performed on a GitHub Codespace with 4 CPU cores and 16 GB RAM. This campaign supports the detailed validation, scalability, AD-overhead, thread-scaling, memory-pressure, and workload-sensitivity analysis. Second, a final cloud campaign evaluates cost-performance and profiler quality on Google Cloud instances, with the headline `sha_cloud_profiler_v4` run executed on `n2-standard-4` in `eu-west1-b`. Claims in the results chapter are limited to these measured local and cloud experiments.

The experimental design follows three principles.

1. **Like-for-like workloads:** every engine is evaluated on the same workload configuration, random seed policy, path count, model parameters, and timing procedure.
2. **Separation of concerns:** performance differences are isolated by varying one main factor at a time, such as engine, AD mode, path count, time-step count, or requested thread count.
3. **Auditable measurement:** each benchmark run stores timing statistics, numerical outputs, memory measurements, software versions, hardware metadata, and the full workload configuration in the shared SQLite database.

Three workloads are used across the complete evaluation. The European GBM workload is used as a simple analytical validation case and low-cost baseline. The parametric local-volatility workload is used as the main local performance workload because it contains time stepping, a nonlinear volatility surface, and a larger computational graph for automatic differentiation. The Asian option is used in the final cloud-profiler experiments as a path-dependent workload that exposes a different backend ranking from the two European-style workloads.

5.2 Connection to the Research Questions

Table 5.1 maps the Chapter 1 research questions directly to the experimental evidence used in Chapter 6.

The main scripts are `run_european_baseline.py`, `run_european_ad_analysis.py`, `run_local_vol_suite.py`, `run_profiler_vs_grid.py`, and `run_guarded_sha_local_vol.py`.

5.3 Execution Environment

The detailed local experiments are executed in a GitHub Codespace with 4 CPU cores and 16 GB RAM. The exact CPU model, operating system, Python version, NumPy version, JAX version, and other software metadata are captured automatically in the benchmark database for every run.

Table 5.1: Mapping from research questions to experimental evidence.

RQ	Evidence collected
RQ1	Runtime, throughput and memory comparisons across NumPy, JAX/XLA, C++/OpenMP and Rust/Rayon using matched European, local-volatility and cloud workloads.
RQ2	Matched JAX no-AD, forward-mode and reverse-mode runs; C++ and Rust are interpreted as compiled no-AD baselines, not C++/Rust AD-library measurements.
RQ3	Sweeps over path count, local-volatility time-step count, payoff type and surface preset to test ranking changes and bottlenecks.
RQ4	Priced Google Cloud CPU runs using measured runtime and marginal compute cost per run from stored instance metadata.
RQ5	Full-grid, old-profiler, plain SHA and guarded SHA comparisons using full runs saved, Pareto recovery, regret and promotion traces.
RQ6	Explicit discussion of validation, uncertainty, RNG differences, warm/cold timing, cloud-cost scope, thread control and generalisation limits.

The use of a single local environment is deliberate: it reduces variation from uncontrolled hardware changes and allows the local experiments to focus on the framework, workload definitions, engines, AD modes, and CPU scaling behaviour. Cloud conclusions are drawn only from the later cloud-profiler experiments and are tied to their recorded instance types, regions, pricing assumptions, and software environment.

5.4 Workloads

5.4.1 European Option under Geometric Brownian Motion

The validation workload prices a European call option under geometric Brownian motion (GBM), using the Black–Scholes model as the analytical reference [22]. Under the risk-neutral measure, the terminal asset price is

$$S_T = S_0 \exp \left[\left(r - \frac{1}{2} \sigma^2 \right) T + \sigma \sqrt{T} Z \right], \quad Z \sim \mathcal{N}(0, 1).$$

The Monte Carlo estimator for the discounted call price is

$$\hat{P} = e^{-rT} \frac{1}{M} \sum_{i=1}^M \max(S_T^{(i)} - K, 0),$$

where M is the number of simulated paths. Unless otherwise stated, the fixed parameters are

$$S_0 = 100, \quad K = 100, \quad r = 0.05, \quad \sigma = 0.20, \quad T = 1.0, \quad \text{seed} = 42.$$

This workload is used for analytical validation, simple baseline performance, path-count scalability, and JAX AD correctness. Its main advantage is that the Black–Scholes analytical price and Greeks provide a direct reference for the Monte Carlo price, Delta, Vega, and Rho.

5.4.2 European Option under Parametric Local Volatility

The main performance workload prices the same European payoff, but replaces constant volatility with a parametric local-volatility surface. The asset path is advanced using a log-Euler discretisation over N time steps:

$$S_{t+\Delta t} = S_t \exp \left[\left(r - \frac{1}{2} \sigma(S_t, t)^2 \right) \Delta t + \sigma(S_t, t) \sqrt{\Delta t} Z_t \right], \quad Z_t \sim \mathcal{N}(0, 1).$$

The local volatility surface is parameterised as

$$\sigma(S, t) = \sigma_{\min} + \text{softplus} \left(a_0 + a_1 \log(S/S_0) + a_2 \log(S/S_0)^2 + b_1 t \right).$$

The softplus function is

$$\text{softplus}(x) = \log(1 + e^x).$$

In implementation, softplus is evaluated using the numerically stable form

$$\text{softplus}(x) = \max(x, 0) + \log(1 + e^{-|x|}),$$

which reduces overflow risk for large positive inputs.

The default local-volatility setting uses daily time steps,

$$N = 252, \quad \sigma_{\min} = 0.01, \quad a_1 = -0.10, \quad a_2 = 0.20, \quad b_1 = 0.00.$$

The parameter a_0 is chosen so that the at-the-money time-zero volatility equals 0.20:

$$a_0 = \log(\exp(0.20 - \sigma_{\min}) - 1).$$

This local-volatility workload is more demanding than the single-step GBM case because each path requires repeated state updates and repeated evaluation of a nonlinear volatility function. It is therefore used as the main workload for engine comparison, AD overhead analysis, path-count scaling, time-step scaling, thread scaling, surface-shape sensitivity, and local resource-usage analysis.

5.4.3 Asian Arithmetic-Average Option

The final cloud-profiler experiments also include an Asian call option whose payoff depends on the arithmetic average of the simulated asset path. In simplified form, the discounted payoff estimator is

$$\hat{P}_{\text{Asian}} = e^{-rT} \frac{1}{M} \sum_{i=1}^M \max\left(\frac{1}{N} \sum_{j=1}^N S_{t_j}^{(i)} - K, 0\right).$$

This workload is included because it is path dependent: the payoff depends on the trajectory rather than only on the terminal value. It therefore introduces a different memory-access and compilation pattern from the European GBM baseline. In the final cloud comparison, the Asian workload is evaluated for supported Python/JAX configurations and contributes to the AD-overhead, cost-performance, and profiler-selection results.

5.5 Engines and Software Stacks

Table 5.2 summarises the engine configurations used in the final experiments. Engines that are not available in a particular run are skipped and recorded as unavailable rather than silently excluded.

The exact Python, NumPy, JAX, compiler, operating-system, CPU, memory, and available backend metadata are captured per run and stored in the database. This avoids relying on manually written machine descriptions, which can become incomplete or stale.

5.6 Experimental Variables

The experiments are designed around a small set of controlled variables.

5.6.1 Independent Variables

The main independent variables are

- **Engine:** NumPy CPU, JAX/XLA, C++/OpenMP, and Rust/Rayon.
- **Path count:** the number of Monte Carlo paths, denoted by M .
- **Time steps:** the number of simulated time steps, denoted by N . For the exact GBM workload, $N = 1$ implicitly because the terminal distribution is sampled directly. For local volatility, N is explicitly swept in selected experiments.
- **AD mode:** no AD, forward-mode AD, or reverse-mode AD. Forward and reverse modes are evaluated through the JAX engine.
- **Thread count:** requested CPU worker count for strong and weak thread-scaling experiments.

Table 5.2: Engine configurations used in the local experimental programme.

Engine label	Language	Backend	Role in evaluation
cpu	Python	NumPy / CPU	Reference vectorised Python implementation. Used as a baseline software stack for no-AD CPU throughput and correctness checks.
jax	Python	JAX / XLA	Compiled AD-capable implementation. Used for no-AD performance, forward-mode AD, reverse-mode AD, and XLA compiled execution behaviour.
cpp	C++	OpenMP	Compiled CPU implementation. Used for low-level CPU throughput and thread-scaling comparison when the extension module is built successfully.
rust	Rust	Rayon	Compiled CPU implementation. Used for language/runtime comparison and parallel CPU throughput on the local-volatility workload.

- **Volatility-surface preset:** local-volatility surface shape, such as flat, equity-skew, or strong-smile, where supported by the experiment runner.
- **Cloud instance:** Google Cloud instance type and region for final cost-performance and profiler experiments.
- **Profiler policy:** full grid, old probe-based profiler, or successive halving.

5.6.2 Dependent Variables

The main dependent variables are

- mean runtime in milliseconds;
- runtime standard deviation, minimum, and maximum;
- throughput in paths per second;
- peak process memory usage in megabytes;
- Monte Carlo price estimate;
- Greek estimates for AD-enabled GBM experiments;
- absolute and relative error against analytical references where available;
- AD overhead ratio;
- requested and observed process thread counts where available;
- coefficient of variation for selected stability experiments.

5.6.3 Controlled Variables

Unless explicitly swept, the following are held fixed:

- option parameters S_0 , K , r , σ , and T ;
- random seed;
- workload definition;
- local-volatility surface parameters;

- warmup count and timed repetition count within each experiment;
- database schema and result-writing path;
- timing method;
- correctness thresholds used for validation.

5.7 Metrics

For each benchmark cell, the script performs a configurable number of warmup runs followed by a configurable number of timed runs. Warmup timings are discarded. The timed runs produce

$$\bar{t}_{\text{ms}} = \frac{1}{R} \sum_{j=1}^R t_j,$$

where R is the number of timed repetitions and t_j is the runtime of the j th repetition in milliseconds.

Throughput is reported as

$$\text{throughput} = \frac{M}{\bar{t}_{\text{ms}} \times 10^{-3}} \quad \text{paths per second.}$$

For AD experiments, the overhead ratio is computed against a matched no-AD baseline:

$$\text{AD overhead} = \frac{\bar{t}_{\text{AD}}}{\bar{t}_{\text{none}}},$$

where both runs use the same engine, workload, path count, time-step count, warmup count, and number of timed repetitions.

For the GBM workload, numerical error is measured against the analytical Black–Scholes reference:

$$\text{relative error} = \frac{|x_{\text{MC}} - x_{\text{ref}}|}{\max(|x_{\text{ref}}|, \epsilon)},$$

where x may be the price, Delta, Vega, or Rho, and ϵ is a small constant used only to avoid division by zero.

For the local-volatility workload, no closed-form analytical reference is available for the chosen parametric surface. Validation is therefore against independent implementations of the same discretised model, not against a closed-form financial reference. This establishes cross-implementation consistency for the chosen discretisation where engines agree, but does not prove correctness of the continuous-time model or the financial modelling assumptions. Where available, the Monte Carlo standard error is used to interpret whether differences between engines are consistent with sampling variation.

For selected stability experiments, the coefficient of variation is reported as

$$\text{CV} = \frac{s_t}{\bar{t}_{\text{ms}}},$$

where s_t is the sample standard deviation of the timed repetitions.

5.8 Measurement Methodology

5.8.1 Warmup and JIT Compilation

Each script performs warmup repetitions before measurement. This is especially important for JAX/XLA, where the first invocation may include tracing and compilation. Warm timings measure repeated-production performance after compilation has occurred. They are appropriate for long-running or repeated workloads, but they understate latency and cost for one-off invocations. Cold-start JAX timings are therefore discussed separately where relevant rather than mixed into the steady-state tables.

This means the cloud cost-performance results should be interpreted as steady-state cost-performance. That is the relevant model for repeated simulation or service-style deployment, but not for one-off serverless or interactive jobs where JAX/XLA compilation latency would be paid each time. A complete cold-start study would report both first-call and warm-call costs for each JAX workload and is treated as extension work rather than part of the final measured evidence.

5.8.2 Cloud Cost Model

Cloud cost is estimated as marginal compute cost: measured runtime multiplied by the published or configured hourly instance price. These estimates exclude VM startup and shutdown time, idle time, storage, orchestration overhead, retries, network costs, and any billing granularity effects. They are therefore suitable for comparing relative cost-performance between configurations under the benchmark protocol, but should not be interpreted as exact end-to-end cloud bills.

5.8.3 Repeated Measurements

Each benchmark cell is measured over multiple timed repetitions. The database stores the mean, standard deviation, minimum, and maximum runtime rather than a single timing. This makes it possible to identify unstable measurements and avoids drawing conclusions from one-off timing noise.

Most main experiments use 7 timed runs and 3 warmup runs. Smoke tests use 1 timed run and 1 warmup run and are not treated as reported performance results. Selected stability experiments use a larger number of timed runs to estimate measurement variability.

5.8.4 Randomness and Correctness

All workloads use an explicit seed. Identical seeds improve reproducibility within each engine, but bit-identical output is not required across engines because different engines may use different random-number generators. Cross-engine correctness is therefore assessed using analytical references where available, and by checking that estimates converge within reasonable Monte Carlo error as M increases.

For the European GBM workload, correctness is assessed directly against the Black-Scholes analytical price and Greeks. For local volatility, the emphasis is on cross-engine consistency and convergence with increasing path count, since the chosen parametric local-volatility model does not have the same closed-form reference.

5.8.5 Fairness Controls

The following controls are applied across experiments:

- Engines receive the same workload parameters and path counts.
- AD overhead is computed only against matched no-AD runs.
- JAX timings are synchronised so that asynchronous execution does not under-report runtime.
- Thread-scaling cells are run in isolated subprocesses where possible to reduce contamination from persistent thread-pool state.
- Failed, skipped, or unavailable engines are recorded explicitly rather than removed silently from the analysis.
- Smoke-test results are used only to validate execution and are not interpreted as performance results.

5.8.6 Data Exclusion Policy

Runs are excluded from the final comparative analysis only under predefined conditions:

- the run terminates with an exception;
- the detected engine or backend does not match the intended experiment cell;
- the result violates the experiment's correctness threshold;

- required metadata such as engine, workload, path count, time-step count, or runtime is missing;
- the machine is known to have been under unrelated heavy load during the run.

Excluded runs are not deleted from the database. They remain available for audit, but are filtered out of the final completed-run queries.

5.9 Experiment Scripts

Table 5.3 summarises the core experiment scripts used in the final evaluation.

Table 5.3: Core experiment scripts and their roles.

Script	Experiment type or labels	Purpose
<code>run_european_baseline.py</code>	<code>european_baseline</code>	Compares no-AD throughput across available engines on the single-step European GBM workload.
<code>run_european_ad_analysis.py</code>	<code>european_ad_analysis</code>	Measures JAX forward- and reverse-mode AD overhead and validates Greek accuracy against Black–Scholes references.
<code>run_european_scalability.py</code>	<code>european_scalability</code>	Sweeps path count M to measure how runtime and throughput scale on the simple GBM workload.
<code>run_thread_scalability.py</code>	<code>thread_scalability_strong; thread_scalability_weak</code>	Measures strong and weak CPU thread scaling on the GBM workload.
<code>run_localvol_suite.py</code>	<code>LV0–LV10</code>	Runs the local-volatility benchmark suite, including engine comparison, AD-overhead scaling, path-count scaling, time-step scaling, thread scaling, surface-shape sensitivity, timing stability, and local stress testing.
<code>run_profiler_vs_grid.py</code>	<code>sha_cloud_profiler_v4</code>	Runs the final cloud full-grid, old-profiler, and SHA comparison, recording cloud metadata, cost estimates, Pareto-frontier recovery, regret, and promotion decisions.

5.10 Individual Experiment Designs

Table 5.4 summarises the individual experiment designs. The detailed result values are reported in Chapter 6; this section records the purpose, sweep variables and primary outputs used to produce those results without repeating the same script/workload fields in prose.

The local-volatility suite is the main experimental programme because it tests whether conclusions from the simple GBM workload still hold for a multi-step, state-dependent MC+AD workload. The GBM experiments remain important as analytical correctness anchors and as cheaper baselines for interpreting fixed runtime overheads, thread behaviour and JAX AD accuracy.

5.11 Final Cloud and Profiler Experiments

The final experimental extension evaluates cloud cost-performance and the budget-aware profiler. It does not replace the earlier local experiments; those experiments provide the detailed workload, AD, scaling and resource analysis. The cloud experiment instead asks whether the completed framework can use stored metadata and cheap probes to make cost-aware full-benchmark selection decisions.

Table 5.4: Summary of local experiment designs.

Experiment	Script / scope	Purpose and sweep	Primary outputs
European GBM baseline	<code>run_european_baseline.py</code> ; NumPy, JAX, C++, Rust where available	No-AD correctness and low-cost engine comparison against Black–Scholes for a fixed European call.	Runtime, throughput, price, relative price error, memory, environment metadata.
European GBM AD	<code>run_european_ad_analyses.py</code> ; JAX	Compares no AD, forward AD and reverse AD for $M \in \{1k, 5k, 10k, 50k, 100k\}$, using analytical Greeks as references.	AD overhead, Delta/Vega/Rho estimates and errors, runtime, memory.
European scaling	<code>run_european_scalability.py</code> ; available engines	Sweeps $M \in \{1k, 5k, 10k, 50k, 100k\}$ to identify fixed overheads and large- M throughput on the simple GBM workload.	Runtime, throughput, price error and memory as functions of M .
European thread scaling	<code>run_thread_scalability.py</code> ; JAX, C++, Rust where available	Runs strong and weak scaling with requested thread counts 1, 2, 4, distinguishing fixed-work and work-per-thread regimes.	Runtime, throughput, requested/observed threads and scaling efficiency.
Local-volatility suite	<code>run_localvol_suite.py</code> ; LV-0–LV-10	Main performance campaign: smoke test, engine scaling, JAX AD scaling with M , time-step scaling, memory pressure, thread scaling, surface sensitivity, timing stability and stress limits.	Runtime, throughput, price, AD overhead, memory, thread metadata and coefficient of variation.

Table 5.5: Final cloud/profiler experiment design.

Component	Design	Evaluation criteria
Full-grid baseline	<code>sha_cloud_profiler_v4</code> , git commit 13789b8, Google Cloud <code>n2-standard-4</code> in <code>europa-west1-b</code> . The reference grid contains 16 supported candidate configurations and $M \in \{10k, 50k, 100k\}$, giving 48 full benchmark cells across European, Asian and local-volatility workloads.	Reference runtime, cost per run, correctness, AD overhead and Pareto frontier.
Old profiler	Cheap probe runs rank candidate configurations and select a conservative subset; in the final run it selects 30 of the 48 full-grid cells.	Full runs saved, Pareto recovery, runtime/cost regret and rank correlation.
Guarded SHA	The final local-volatility profiler evaluates price-only and AD-required candidate pools separately across <code>e2</code> , <code>n2</code> , <code>n2d</code> , <code>c2d</code> and <code>t2d</code> four-vCPU instances. Probe budgets are $M = 1k, 5k$ and $25k$, with a 100k-path full-grid oracle.	Recommendation quality, full runs saved, regret against the oracle, guard decisions, and whether the guarded policy improves the plain SHA robustness story.

5.12 Execution Order and Analysis Queries

Experiments are run from cheapest and most analytically checkable to most expensive: European GBM validation; European AD validation; GBM path/thread scaling; local-volatility smoke testing; LV-1/LV-2 engine and AD sweeps; LV-3/LV-4 *N*-scaling; LV-5–LV-10 resource, thread, surface, stability and stress tests; then `sha_cloud_profiler_v4`; and finally the guarded local-volatility SHA run across the five cloud instance types. This order keeps GBM as the correctness anchor while making local volatility the main performance benchmark and the cloud experiment the main evidence for cost-aware selection.

The final analysis is produced from completed rows in the shared SQLite database. The views used for Chapter 6 filter by experiment type, engine, AD mode, workload size, status and thread metadata to produce: European baseline throughput; European AD overhead and Greek accuracy; local-volatility no-AD engine scaling; local-volatility JAX AD overhead; strong/weak thread scaling; and cloud profiler decision-quality summaries. The full database and generated CSV artefacts are submitted with the project repository, so the exact predicates can be audited without printing pages of SQL in the main report.

5.13 Reproducibility

Each completed run stores the full workload configuration and environment metadata needed to reproduce or audit the measurement. The database record includes

- the serialised workload configuration;
- a configuration hash where available;
- engine, language, backend, and AD mode;
- path count, time-step count, option parameters, and seed;
- local-volatility surface parameters or preset;
- timing statistics;
- price and Greek outputs where available;
- analytical reference values for GBM runs;
- numerical error metrics where a reference is available;
- memory usage;
- requested and observed thread counts where available;
- CPU, operating-system, Python, NumPy, and JAX metadata.

This metadata ensures that every reported number can be traced back to a specific workload, software stack, and hardware environment. The experiment database is append-only during collection: failed, skipped, and completed runs are all preserved so that the final analysis does not depend on unrecorded manual filtering.

5.14 Summary

The experimental setup defines a controlled methodology for MC+AD workloads, combining a 4-core, 16 GB RAM Codespace campaign with final Google Cloud profiler experiments. The European GBM workload provides analytical validation and a simple baseline. The parametric local-volatility workload is the main detailed performance benchmark, adding multi-step simulation, nonlinear model structure, and higher AD complexity. The Asian workload adds a path-dependent case to the cloud-profiler grid.

Together, the experiments evaluate runtime performance, scalability, automatic differentiation overhead, numerical consistency, thread behaviour, local resource usage, cloud cost-performance, and budget-aware profiler quality across NumPy, JAX/XLA, C++/OpenMP, and Rust/Rayon where supported. The next chapter reports the measurements obtained from these experiments and discusses how far they answer the project research questions.

Chapter 6

Results and Discussion

6.1 Overview

This chapter reports and interprets the benchmark results obtained from the experimental setup in Chapter 5. The European GBM workload is used as a validation and calibration workload: it verifies that the framework produces credible Monte Carlo prices and automatic differentiation outputs against a known Black–Scholes reference. The main performance analysis is then carried out on the parametric local-volatility workload, which is more representative of the target MC+AD setting because it includes time stepping, a nonlinear volatility surface, and a larger differentiation graph.

The results combine two layers of evidence. The first layer is a detailed local CPU study on the Codespace environment, comparing NumPy, JAX/XLA, C++/OpenMP, and Rust/Rayon where available. This provides the main evidence for validation, workload scaling, AD overhead, memory pressure, and thread behaviour. The second layer is the final cloud-profiler study, which adds Google Cloud cost-performance measurements, Pareto-frontier analysis, and old-profiler versus SHA selection quality. The discussion below only draws claims from these measured experiments.

Table 6.1: Result sections and their corresponding evidence.

Section	Experiment source	Main evidence
European GBM validation	European baseline and AD analysis	Black–Scholes price and Greek agreement used to validate the framework.
Local-volatility engine scaling	LV-1	No-AD runtime scaling across NumPy, JAX, C++, and Rust as path count increases.
Local-volatility AD overhead	LV-2 and LV-4	Forward- and reverse-mode JAX AD overhead as a function of path count M and time-step count N .
Time-step scalability	LV-3	Runtime scaling with the number of local-volatility time steps.
Memory and stress limits	LV-5 and LV-10	Resource limits for JAX AD, including observed out-of-memory failures.
Thread scalability	LV-6 and LV-7	Strong and weak scaling for C++/OpenMP, Rust/Rayon, and JAX/XLA.
Surface-shape sensitivity	LV-8	Runtime effect of changing the local-volatility surface preset.
Timing stability	LV-9	Coefficient of variation over 20 repeated runs for key benchmark cells.
Final cloud cost-performance	<code>sha_cloud_profiler_v4</code> and generated cloud artefacts	Runtime, estimated cost per run, and best no-AD configurations for priced cloud instances.
Profiler selection quality	Full-grid, old-profiler, and SHA comparison	Full runs saved, Pareto recovery, runtime regret, cost regret, rank correlation, and SHA promotion behaviour.

6.2 European GBM Validation

The European GBM workload is not the main performance workload in this chapter. Its role is to provide a simple analytical reference before interpreting the more expensive local-volatility experiments. For the standard parameter set $S_0 = K = 100$, $r = 0.05$, $\sigma = 0.20$, and $T = 1$, the Black–Scholes analytical call price [22] is 10.4506, with Delta 0.6368, Vega 37.524, and Rho 53.232.

Table 6.2: European GBM baseline validation at $M = 10,000$.

Engine	Price	Mean ms	Std ms	Paths/s	Rel. price error
NumPy/CPU	10.3452	0.220	0.016	45.5M	1.01%
JAX/XLA	10.2098	0.559	0.051	17.9M	2.30%
C++/OpenMP	10.4658	0.098	0.009	101.6M	0.15%

The baseline GBM results show that all completed engines produce prices close to the analytical Black–Scholes value. The C++ engine is the fastest on this small single-step workload, achieving 101.6 million paths/s, while NumPy reaches 45.5 million paths/s and JAX reaches 17.9 million paths/s. This result is useful as a sanity check, but it is not treated as the main performance story because the workload is too small and too simple to represent the multi-step MC+AD case.

The JAX AD validation also produced accurate Greeks at larger path counts. At $M = 100,000$, forward- and reverse-mode AD gave the same reported Greek estimates to the displayed precision. Delta error was 0.000428 (0.07%), Vega error was 0.259 (0.69%), and Rho error was 0.080 (0.15%). This confirms that the JAX AD implementation is numerically credible before it is used for the local-volatility AD experiments.

6.3 Local-Volatility Smoke Test

LV-0 is a smoke test and is not used as a main performance result. It checks that each available engine can execute the local-volatility workload and write results to the shared database. The test uses only one timed run at $M = 1,000$, so the timings are not statistically meaningful.

Table 6.3: LV-0: smoke test ($M = 1,000$, $N = 252$, one timed run).

Engine	AD mode	Latency ms	Price	Delta / Vega
cpu	none	16.000	10.61771	—
cpp	none	7.427	10.79929	—
rust	none	5.095	12.19688	—
jax	none	10.205	10.33823	—
jax	forward	45.499	10.33823	$\delta = 0.6294$, $\nu = 37.49$
jax	reverse	23.859	10.33823	$\delta = 0.6294$, $\nu = 37.49$

The smoke test confirms that the local-volatility pipeline runs across all four engines. The Rust price is visibly different from the other engines at this very small path count, but this is not interpreted as a correctness failure because $M = 1,000$ is too small for a reliable local-volatility cross-engine comparison and different engines may use different random number generators. The larger M -sweeps below are used for performance interpretation.

Because the local-volatility model has no closed-form price, an additional uncertainty check was run using an independent batched NumPy simulation with $M = 1,000,000$, $N = 252$, and seed 12345. The discounted payoff standard deviation from that run was 14.64743, giving a Monte Carlo standard error of 0.01465 for the reference price. Table 6.4 compares this high- M reference with the largest LV-1 prices. The z -scores use the combined standard error of the reference and the corresponding $M = 500,000$ estimate, using the reference payoff standard deviation as a common estimate of payoff variance.

This check strengthens the numerical interpretation of the local-volatility results. NumPy, C++ and JAX are within approximately 1.3 combined standard errors of the independent reference, so their price differences are consistent with Monte Carlo sampling error. The Rust local-volatility price is outside this range; the discrepancy also persists in the $M = 1,000,000$ stress row. For that reason, Rust local-volatility timings are retained as useful runtime and scaling measurements,

Table 6.4: Local-volatility uncertainty check against an independent batched NumPy reference.

Estimate	M	Price	Difference	z -score
Independent NumPy reference	1,000,000	10.44452 ± 0.01465	—	—
cpu LV-1	500,000	10.47542	+0.03089	1.22
cpp LV-1	500,000	10.46301	+0.01848	0.73
jax LV-1	500,000	10.46082	+0.01629	0.64
rust LV-1	500,000	10.54431	+0.09979	3.93

but the Rust local-volatility price is not used as numerical validation evidence. The most likely explanation is the interaction between Rayon work splitting and per-thread random-number streams, which should be audited before treating Rust as a trusted local-volatility pricing backend.

6.4 Local-Volatility Engine Scalability

LV-1 is the main no-AD engine comparison on the local-volatility workload. The workload uses $N = 252$ time steps and sweeps path count from 10,000 to 500,000.

Table 6.5: LV-1: engine scalability on local volatility ($N = 252$, no AD).

Engine	$M = 10k$ ms	$M = 25k$ ms	$M = 50k$ ms	$M = 100k$ ms	$M = 250k$ ms	$M = 500k$ ms
cpu	100.7 ± 3.5	255.8 ± 7.6	501.5 ± 10.4	995.4 ± 24.7	2448.3 ± 25.7	5022.4 ± 52.0
cpp	69.8 ± 12.9	158.8 ± 19.2	338.6 ± 45.4	633.4 ± 41.8	1621.0 ± 61.9	3062.6 ± 82.1
rust	57.3 ± 9.4	122.1 ± 11.9	239.3 ± 14.1	497.1 ± 32.2	1227.7 ± 55.0	2439.6 ± 57.3
jax	93.4 ± 12.6	199.6 ± 41.4	339.3 ± 47.8	629.9 ± 24.2	1558.3 ± 41.0	3103.7 ± 45.0

The local-volatility workload shows approximately linear runtime growth with path count for all engines. This is expected because the workload cost is roughly proportional to $M \times N$, and N is fixed at 252 in this experiment.

Rust/Rayon is the fastest no-AD engine by measured runtime across the full path-count sweep. At $M = 500,000$, Rust completes the workload in 2439.6 ms, compared with 3062.6 ms for C++, 3103.7 ms for JAX, and 5022.4 ms for NumPy. This corresponds to approximately 0.205 million paths/s for Rust, 0.163 million paths/s for C++, 0.161 million paths/s for JAX, and 0.100 million paths/s for NumPy. Thus, on this local-volatility workload, Rust is approximately $2.1\times$ faster than NumPy and about $1.25\times$ faster than JAX or C++ at the largest tested path count.

The runtime ranking should be read together with the uncertainty check in Table 6.4. The CPU, C++ and JAX prices are statistically consistent with the independent high- M reference, while the Rust local-volatility price requires further RNG-stream validation. The Rust local-volatility timings are reported as engineering performance measurements, but the corresponding prices show a discrepancy relative to the NumPy/JAX/C++ agreement band. Consequently, Rust local-volatility timings should not be interpreted as evidence of a superior validated pricing engine. They indicate the performance of the current implementation under the implemented workload, while correctness requires further investigation.

A notable difference from the simple European GBM workload is that JAX becomes competitive with C++ on the more expensive multi-step workload. This suggests that the fixed overheads of XLA execution are less important once each benchmark cell contains enough path-level and time-step work.

6.5 Local-Volatility AD Scaling with Path Count

LV-2 measures the cost of JAX forward- and reverse-mode AD on the local-volatility workload as the number of paths increases. All runs use $N = 252$.

The main AD result is clear: reverse-mode AD is substantially cheaper than forward-mode AD for this scalar-output local-volatility pricing workload. Across the tested path counts, forward-mode overhead is roughly $4.1\text{--}6.0\times$, while reverse-mode overhead is roughly $2.1\text{--}2.4\times$. This is consistent

Table 6.6: LV-2: JAX AD path-count scaling on local volatility ($N = 252$).

M	None ms	Forward ms	Reverse ms	Forward OH	Reverse OH
5k	42.4 ± 0.2	210.4 ± 4.6	90.3 ± 4.9	$4.94\times$	$2.13\times$
10k	98.8 ± 19.8	437.0 ± 7.6	179.4 ± 5.2	$5.02\times$	$2.08\times$
25k	166.1 ± 20.7	985.7 ± 51.2	373.8 ± 14.3	$6.00\times$	$2.35\times$
50k	340.9 ± 40.9	1511.3 ± 31.9	654.6 ± 20.5	$4.87\times$	$2.10\times$
100k	644.1 ± 25.8	3121.3 ± 106.0	1312.0 ± 35.2	$4.86\times$	$2.06\times$
250k	1558.2 ± 24.0	6252.4 ± 109.7	3627.0 ± 82.6	$4.08\times$	$2.34\times$

with the expected AD trade-off: the pricing function has many input parameters but produces a scalar price, making reverse mode a more suitable differentiation strategy.

The overhead does not vanish at large M . At $M = 250,000$, reverse mode still costs $2.34\times$ the no-AD runtime, while forward mode costs $4.08\times$. This shows that AD remains a first-order performance concern for MC workloads even when the underlying simulation is large enough to amortise JAX/XLA dispatch overheads.

6.6 Local-Volatility Scaling with Time Steps

LV-3 separates path-count scaling from time-discretisation scaling by fixing $M = 50,000$ and varying the number of local-volatility time steps.

Table 6.7: LV-3: time-step scaling on local volatility ($M = 50,000$, no AD).

Engine	$N = 12$	$N = 52$	$N = 126$	$N = 252$	$N = 504$
cpp	17.6 ± 5.9	66.7 ± 8.3	170.2 ± 25.7	328.0 ± 49.3	633.2 ± 32.2
rust	12.8 ± 3.1	49.9 ± 3.9	132.3 ± 7.5	256.1 ± 18.6	519.4 ± 63.0
jax	20.5 ± 3.8	61.7 ± 4.8	171.4 ± 27.3	324.0 ± 40.8	648.7 ± 39.5

Runtime increases approximately linearly with N , which confirms that the local-volatility workload is dominated by repeated time-step updates rather than a fixed one-off overhead. For example, Rust increases from 256.1 ms at $N = 252$ to 519.4 ms at $N = 504$, close to a $2\times$ increase when the number of time steps doubles. JAX and C++ show the same general pattern.

Rust remains the fastest engine by runtime in this experiment, followed by C++ and JAX. The difference between C++ and JAX is small at $N = 252$, but both are slower than Rust. This reinforces the LV-1 result that the Rust/Rayon implementation is the strongest no-AD local-volatility runtime engine in the current local environment.

6.7 Local-Volatility AD Scaling with Time Steps

LV-4 measures how JAX AD overhead changes with the number of time steps. The path count is fixed at $M = 25,000$.

Table 6.8: LV-4: JAX AD time-step scaling on local volatility ($M = 25,000$).

N	None ms	Forward ms	Reverse ms	Forward OH	Reverse OH
12	7.5 ± 0.7	50.1 ± 3.5	18.6 ± 0.8	$7.14\times$	$2.47\times$
52	28.7 ± 0.9	203.3 ± 3.8	73.9 ± 2.2	$5.11\times$	$2.51\times$
126	72.6 ± 2.0	504.7 ± 27.0	186.9 ± 9.8	$6.81\times$	$2.62\times$
252	156.7 ± 5.3	974.6 ± 15.8	372.4 ± 14.4	$6.62\times$	$2.52\times$
504	329.5 ± 27.7	1901.5 ± 13.5	738.4 ± 14.7	$6.32\times$	$2.33\times$

The AD overhead pattern is stable across simulation length. Reverse mode remains close to $2.3\text{--}2.6\times$ the no-AD runtime, while forward mode remains much more expensive, usually above $5\times$. Increasing the number of time steps therefore increases absolute runtime for all modes, but does not change the conclusion that reverse mode is the better AD strategy for this scalar-output workload.

The $N = 504$ case is particularly important because it doubles the default daily time-step count. Even there, reverse mode completes in 738.4 ms, compared with 1901.5 ms for forward mode. This gives a strong practical argument for reverse-mode AD when many sensitivities are required from a single price output.

6.8 Memory Pressure and Practical Resource Limits

Memory behaviour is analysed separately because runtime alone does not show whether a method remains usable at larger M and N . This matters especially for reverse-mode AD, where intermediate values from the scan may need to be retained or reconstructed during the backward pass.

Table 6.9: LV-5: JAX memory-pressure experiment.

N	AD mode	$M = 25k$ ms	$M = 50k$ ms	$M = 100k$ ms	$M = 250k$ ms
252	none	152.4 ± 5.7	315.2 ± 23.2	650.8 ± 38.4	1527.9 ± 20.5
	reverse	380.1 ± 21.1	649.4 ± 22.2	1332.9 ± 64.3	3573.5 ± 51.4
	forward	961.8 ± 16.9	1515.7 ± 38.3	3103.6 ± 62.7	6278.7 ± 161.8
504	none	334.5 ± 61.9	647.6 ± 32.6	1239.6 ± 24.9	3102.3 ± 58.6
	reverse	754.8 ± 12.3	1310.9 ± 26.0	2595.1 ± 78.6	OOM
	forward	1979.6 ± 100.8	3006.0 ± 49.3	6101.9 ± 133.7	12392.9 ± 121.3

OOM: the JAX reverse-mode cell at $M = 250k$, $N = 504$ was killed by the kernel out-of-memory mechanism.

The most important memory result is not the small process-level RSS deltas, which are noisy under JAX/XLA, but the observed out-of-memory boundary. JAX reverse mode succeeds at $M = 100k$, $N = 504$, but fails at $M = 250k$, $N = 504$. It also succeeds at $M = 500k$, $N = 252$ in the stress test below, but fails at $M = 1M$, $N = 252$. These failures provide direct evidence that reverse-mode AD has a practical memory envelope on the 4-core, 16 GB Codespace environment.

Forward mode is slower than reverse mode, but it completes all LV-5 cells, including $M = 250k$, $N = 504$. This gives a useful trade-off: reverse mode is substantially faster for moderate problem sizes, but forward mode may be more robust in some large memory-stressed settings.

6.9 Thread Scalability on Local Volatility

LV-6 and LV-7 evaluate CPU thread scaling on the local-volatility workload. The experiments use requested thread counts $T = 1, 2, 4$. C++ and Rust use explicit OpenMP/Rayon parallelism, while JAX uses XLA’s own CPU execution runtime and is not expected to respond directly to `OMP_NUM_THREADS`.

6.9.1 Strong Scaling

Table 6.10: LV-6: strong thread scaling ($M = 100,000$, $N = 252$).

Engine	$T = 1$ ms	$T = 2$ ms	$T = 4$ ms	S_2	S_4	Efficiency ₄
cpp	1944.4 ± 37.3	973.9 ± 5.9	629.9 ± 31.7	$2.00\times$	$3.09\times$	77.2%
rust	1787.3 ± 2.7	900.0 ± 4.1	488.1 ± 34.6	$1.99\times$	$3.66\times$	91.5%
jax	642.4 ± 20.7	637.5 ± 32.3	642.7 ± 19.4	$1.01\times$	$1.00\times$	25.0%

$S_k = T_1/T_k$. JAX uses XLA’s own CPU parallelism and does not show meaningful sensitivity to the requested OpenMP thread count.

C++ and Rust both scale strongly with requested thread count. Rust improves from 1787.3 ms at one thread to 488.1 ms at four threads, giving a $3.66\times$ speedup and 91.5% four-thread efficiency. C++ improves from 1944.4 ms to 629.9 ms, giving a $3.09\times$ speedup and 77.2% efficiency. This shows that the local-volatility workload exposes substantial path-level parallelism to compiled CPU engines.

JAX does not scale with the requested thread count in this experiment. Its runtime remains close to 640 ms for $T = 1, 2, 4$. This does not mean JAX is single-threaded; rather, it suggests that

XLA manages CPU parallelism internally and is not controlled by the same OpenMP/Rayon knobs used by C++ and Rust.

6.9.2 Weak Scaling

Table 6.11: LV-7: weak thread scaling (base $M = 50,000$ paths/thread, $N = 252$).

Engine	$T = 1, M = 50k$ ms	$T = 2, M = 100k$ ms	$T = 4, M = 200k$ ms	TP speedup
cpp	961.8 ± 7.6	975.2 ± 13.8	1238.7 ± 29.6	$3.11\times$
rust	893.6 ± 2.1	913.4 ± 33.7	969.9 ± 34.2	$3.68\times$
jax	316.2 ± 16.1	675.0 ± 95.8	1252.0 ± 32.1	$1.01\times$

Throughput speedup is computed as $(4M_{\text{base}}/T_4)/(M_{\text{base}}/T_1)$. Ideal weak scaling would give $4\times$ throughput at approximately constant latency.

Weak scaling again favours the compiled engines. Rust achieves a $3.68\times$ throughput speedup from one to four requested threads, while C++ achieves $3.11\times$. Latency is not perfectly constant, especially for C++ at $T = 4$, but both compiled engines make effective use of additional requested threads.

JAX shows almost no weak-scaling throughput gain under the requested thread sweep. Its latency increases from 316.2 ms at $M = 50k$ to 1252.0 ms at $M = 200k$, giving only $1.01\times$ throughput speedup. This supports the interpretation that JAX/XLA’s CPU execution strategy is not controlled by the same explicit thread-count mechanism as the compiled OpenMP/Rayon engines.

6.10 Surface-Shape Sensitivity

LV-8 tests whether changing the local-volatility surface shape materially changes runtime. The experiment compares flat, equity-skew, and strong-smile presets at $M = 50,000$, $N = 252$.

Table 6.12: LV-8: surface-shape sensitivity ($M = 50,000$, $N = 252$).

Engine	AD	Flat ms	Equity-skew ms	Strong-smile ms	Reverse OH
jax	none	314.9 ± 26.8	318.1 ± 24.6	325.5 ± 14.8	—
	reverse	669.9 ± 47.3	663.4 ± 27.4	682.4 ± 77.1	$2.09\text{--}2.15\times$
cpp	none	347.8 ± 64.5	327.9 ± 14.4	341.2 ± 75.3	—

JAX prices were 10.38974, 10.39082, and 10.40619 for the flat, equity-skew, and strong-smile presets respectively. C++ prices were 10.44431, 10.44826, and 10.45843.

Surface shape has little effect on runtime. For JAX no-AD execution, the three presets all complete between 314.9 ms and 325.5 ms. For JAX reverse mode, overhead remains in a narrow range of $2.09\text{--}2.15\times$. This suggests that the cost of the local-volatility workload is dominated by the number of paths, number of time steps, random number generation, and scan structure, rather than by the particular parameter values used in the polynomial volatility surface.

This experiment is still useful because it demonstrates that the framework can vary model structure while keeping the same execution harness. That is important for extensibility: new model presets can be added without changing the benchmarking methodology.

6.11 Timing Stability

LV-9 repeats selected key cells for 20 timed runs and reports the coefficient of variation (CV). This checks whether the benchmark results are stable enough to support comparisons.

The timing stability experiment shows acceptable variability for a shared Codespace environment. CV values range from 3.23% for JAX reverse mode to 6.86% for C++. This is stable enough to support the main engine and AD comparisons, especially because the observed performance differences between engines and AD modes are much larger than the measured run-to-run variation.

Table 6.13: LV-9: timing stability ($M = 100,000$, $N = 252$, 20 timed runs).

Engine	AD	Mean ms	Std ms	Min ms	Max ms	CV
jax	none	629.5	29.9	596.4	722.9	4.75%
jax	reverse	1275.3	41.2	1227.5	1390.2	3.23%
cpp	none	632.7	43.4	594.0	739.4	6.86%
rust	none	486.3	22.2	459.9	529.9	4.56%

$$CV = \sigma/\mu \times 100\%.$$

Rust is both the fastest no-AD engine in this stability subset and one of the more stable engines, with a CV of 4.56%. JAX reverse mode is the most consistent cell, with a CV of 3.23%, likely because the compiled execution path is stable after warmup.

6.12 Stress Limit

LV-10 probes the largest local-volatility workloads attempted on the Codespace environment. It compares the fastest compiled no-AD engines against JAX reverse-mode AD at $N = 252$.

Table 6.14: LV-10: stress limit ($N = 252$, large M).

Engine	AD	$M = 500k$ ms	Mem MB	$M = 1M$ ms
cpp	none	3099.3 ± 41.9	—	6363.0 ± 202.2
rust	none	2422.2 ± 36.7	—	5028.3 ± 36.9
jax	reverse	6878.3 ± 103.9	527.9	OOM

JAX reverse mode at $M = 1M$, $N = 252$ was killed by the kernel out-of-memory mechanism.

The stress test reinforces two runtime conclusions. First, Rust remains faster than C++ at large local-volatility path counts: at $M = 1M$, Rust completes in 5028.3 ms, while C++ takes 6363.0 ms. Second, JAX reverse-mode AD has a practical memory limit in the current environment. It completes $M = 500k$, $N = 252$, but fails at $M = 1M$, $N = 252$.

Together with LV-5, the OOM cases locate the practical JAX reverse-mode boundary between successful $M = 500k$, $N = 252$ and failed $M = 1M$, $N = 252$, and between successful $M = 100k$, $N = 504$ and failed $M = 250k$, $N = 504$. This is one of the most important resource-utilisation findings in the chapter: reverse-mode AD is faster than forward mode for moderate local-volatility workloads, but memory becomes the limiting factor at larger path-count/time-step products.

6.13 Final Cloud Cost-Performance and Profiler Results

The local experiments above explain the detailed behaviour of the framework on controlled CPU workloads. The final cloud experiment extends that evidence by asking whether the framework can make cost-aware benchmark-selection decisions from stored cloud metadata and cheap probes. The headline experiment is `sha_cloud_profiler_v4`, run at git commit `13789b8` on Google Cloud `n2-standard-4` in `eu-west1-b`. The full-grid baseline contains 16 candidate configurations evaluated at three full path counts, for 48 full benchmark cells.

6.13.1 Correctness and AD Overhead in the Final Cloud Run

The European GBM workload is again used as the analytical validation point. Table 6.15 shows the stored JAX AD validation at $M = 100,000$. Forward and reverse mode produce the same reported price and Greeks in this experiment. The price error is 0.357%, and the absolute errors for Delta, Vega, and Rho are small relative to the scale of the reported Greeks.

Table 6.16 reports JAX AD overhead at $M = 100,000$ on `n2-standard-4`. The final cloud results broaden the earlier local-volatility-only AD discussion by including European, Asian, and local-volatility workloads in the same cloud comparison.

Table 6.15: European GBM AD validation against Black–Scholes analytical values at $M = 100,000$.

AD mode	Price err.	Delta err.	Vega err.	Rho err.
forward	0.357%	0.000428	0.259	0.080
reverse	0.357%	0.000428	0.259	0.080

Table 6.16: JAX automatic differentiation overhead at $M = 100,000$ on n2-standard-4.

Workload	AD mode	Runtime ms	Overhead
Asian	reverse	1148.4	$1.81\times$
European	forward	4.783	$1.96\times$
European	reverse	5.208	$2.10\times$
Asian	forward	1552.1	$2.45\times$
Local volatility	reverse	2169.7	$2.55\times$
Local volatility	forward	3768.5	$4.42\times$

The measured overheads range from $1.81\times$ for Asian reverse mode to $4.42\times$ for local-volatility forward mode. The ordering is workload dependent: reverse mode is cheaper than forward mode for the Asian and local-volatility measurements, while the European forward and reverse overheads are closer. This reinforces the earlier conclusion that AD mode should be measured under the actual workload rather than chosen by a fixed rule of thumb.

6.13.2 Cloud Cost-Performance

Table 6.17 summarises the best no-AD cost-performance rows for priced cloud instances at $M = 100,000$. Cloud cost is estimated as marginal compute cost: measured runtime multiplied by the published or configured hourly instance price. These estimates exclude VM startup and shutdown time, idle time, storage, orchestration overhead, retries, network costs, and billing granularity effects. They are therefore suitable for relative cost-performance comparison under the benchmark protocol, not as exact end-to-end cloud bills.

Table 6.17: Best no-AD cost-performance per priced cloud instance at $M = 100,000$.

Workload	Instance	Best engine	Runtime ms	Cost/run USD	M paths/s
Asian	n2-standard-4	jax/none	634.6	$3.42\text{e-}05$	0.158
Asian	t2d-standard-4	jax/none	343.9	$1.62\text{e-}05$	0.291
European	n2-standard-4	rust/none	0.539	$2.91\text{e-}08$	185.495
European	t2d-standard-4	rust/none	0.356	$1.68\text{e-}08$	281.263
Local volatility	n2-standard-4	rust/none	488.9	$2.64\text{e-}05$	0.205
Local volatility	t2d-standard-4	rust/none	443.6	$2.10\text{e-}05$	0.225

Figure 6.1 visualises the same relationship, with cost per run shown on a logarithmic axis because the absolute marginal compute costs are small and differ by workload. In the stored priced comparison, **t2d-standard-4** is faster and lower cost per run than **n2-standard-4** for the listed best no-AD workloads. This is an empirical result for the measured workloads and pricing assumptions, not a claim about all machine types or future cloud prices.

6.13.3 Pareto Frontier

Figure 6.2 plots runtime against estimated cost for the final **n2-standard-4** full-grid comparison. Because the plotted configurations are on the same instance type, runtime and cost are closely linked. The value of the Pareto analysis is therefore not that it reveals a complex two-dimensional trade-off on one machine, but that it gives a precise target for profiler recovery: a profiler should retain the configurations that the full grid identifies as non-dominated under the measured metrics.




Figure 6.1: Best no-AD cloud configurations at $M = 100,000$. The y-axis uses log-scale marginal compute cost per run; lower-left points indicate faster and cheaper measured configurations under the benchmark protocol.




Figure 6.2: Runtime versus estimated cost for the final n2-standard-4 grid. Pareto-optimal points are highlighted.

6.13.4 Old Profiler versus SHA

Table 6.18 is the main final cloud result. The full grid contains 48 full benchmark runs. The old profiler selects 30 runs, saving 37.5%. SHA selects 24 runs, saving 50.0%. Both profilers recover 100.0% of the full-grid Pareto frontier and incur 0.0% runtime and cost regret relative to the full-grid best. SHA also achieves a stored final probe-to-full Spearman rank correlation of 1.000, compared with 0.935 for the old profiler evidence.

Table 6.18: Full-grid, old-profiler and SHA comparison for the final **n2-standard-4** cloud experiment. SHA preserves the full-grid decision on this slice while reducing full-scale runs.

Metric	Full grid	Old profiler	SHA
Candidate configurations	16	16	16
Full benchmark runs	48	30	24
Full runs saved	0	18	24
Percentage saved	0.0%	37.5%	50.0%
Pareto recovery	baseline	100.0%	100.0%
Runtime regret	baseline	+0.0%	+0.0%
Cost regret	baseline	+0.0%	+0.0%
Spearman ρ probe–full	–	0.935	1.000
Probe path-budget saving	–	–	95.5%
Mean scaling-law error	–	–	9.7%

This result is promising for the **n2-standard-4** slice, but it should not be read as proof that SHA is generally reliable across workloads, instance families, pricing snapshots, or noisy probe rankings. In the measured headline grid, SHA halves the number of full-scale cloud benchmark cells while preserving the full-grid decision quality on the reported metrics. The practical significance is that the expensive part of the experiment is not merely shifted elsewhere: the SHA probe path budget is 95.5% lower than evaluating the grid at full scale, and the mean scaling-law extrapolation error is 9.7%.

6.13.5 Guarded SHA Local-Volatility Profiler

The robustness issue above motivated a guarded SHA experiment focused on the local-volatility deployment problem. The candidate space is split into fair task pools: an all-implemented-backends price-only pool, a validated-backend price-only pool that excludes Rust because of the local-volatility price discrepancy discussed in Section 6.15.2, and an AD-required pool comparing only configurations that produce Greeks. The run evaluates five four-vCPU Google Cloud instance types (**e2**, **n2**, **n2d**, **c2d**, and **t2d**) with $M = 1k$, $5k$, and $25k$ probes and a $100k$ -path full-grid oracle. The clean combined artefact contains 120 observations.

Table 6.19: Guarded SHA local-volatility recommendations and pure runtime/cost baselines. The weighted recommendation is compared with the pure fastest and pure cheapest configuration in the same candidate pool.

Pool/task	Objective	Weighted rec.	Fastest	Cheapest	Regret
All price-only	Speed	rust/none/t2d	rust/none/c2d	rust/none/t2d	0.0%
All price-only	Balanced	rust/none/t2d	rust/none/c2d	rust/none/t2d	0.0%
All price-only	Cost	rust/none/t2d	rust/none/c2d	rust/none/t2d	0.0%
Validated price-only	Speed	cpp/none/t2d	cpp/none/t2d	cpp/none/t2d	0.0%
Validated price-only	Balanced	cpp/none/t2d	cpp/none/t2d	cpp/none/t2d	0.0%
Validated price-only	Cost	cpp/none/t2d	cpp/none/t2d	cpp/none/t2d	0.0%
AD-required	Speed	jax/reverse/t2d	jax/reverse/c2d	jax/reverse/t2d	0.0%
AD-required	Balanced	jax/reverse/t2d	jax/reverse/c2d	jax/reverse/t2d	0.0%
AD-required	Cost	jax/reverse/t2d	jax/reverse/c2d	jax/reverse/t2d	0.0%

Table 6.19 shows a stronger and more defensible profiler story than the unguarded SHA audit. The repeated **t2d** selection is not a transcription error: the objective is a weighted runtime–cost score, not pure elapsed time. **c2d** is the fastest pure runtime choice, but **t2d** is only 6.0% slower and 24.3% cheaper for price-only, and 4.9% slower and 25.5% cheaper for AD-required. This is enough for **t2d** to win even the speed-sensitive 0.8/0.2 objective.

Table 6.20: Same-grid profiler comparison for the balanced guarded-SHA objective. Guarded SHA uses more full runs than plain SHA, but adds explicit safeguards.

Method	Used	Saved	Selected config.	Obj.	Run	Cost	Validity status
Full-grid oracle	30	0	Price: rust perf./ cpp valid; AD: jax	0.0%	0.0%	0.0%	Reference grid
Plain SHA	7	23	Price: rust perf./ cpp valid; AD: jax	0.0%	0.0%	0.0%	Less guarded
Guarded SHA	17	13	Price: rust perf./ cpp valid; AD: jax	0.0%	0.0%	0.0%	Preferred policy

The all-backend price-only pool selects **rust/none/t2d**, but this is a performance recommendation only. Because Rust local-volatility pricing has not passed the same cross-engine consistency check as NumPy, C++ and JAX, the validated price-only deployment recommendation is instead **cpp/none/t2d**. The AD-required recommendation is **jax/reverse/t2d**; these local-volatility sensitivities are internally consistent JAX AD outputs rather than closed-form-validated Greeks.

The 0.0% regret values are measured against the observed full-grid oracle for this tested local-volatility configuration, candidate engines, AD modes, instance types and 100k-path target budget. They do not imply the true global optimum over all cloud instances, thread counts, workloads or larger path budgets. Table 6.20 shows that plain SHA also recovers the observed oracle on this grid while selecting fewer full runs; guarded SHA is reported as the safer policy because it trades some of that saving for explicit near-tie, engine-diversity, instance-family and scaling-improvement safeguards. Cloud timing variation means the small **c2d/t2d** gaps should be interpreted as near-tie deployment evidence rather than permanent hardware rankings.

6.13.6 SHA Progression in the Earlier Multi-Workload Grid

This subsection returns to the earlier multi-workload `sha_cloud_profiler_v4` experiment and visualises the promotion path on the **n2-standard-4** grid. Figure 6.3 shows that all 16 configurations are evaluated at the cheapest probe level, then the active set is reduced as the path budget increases. The final selected set contains the configurations promoted for full benchmarking after the profiler’s conservative safeguards are applied.




Figure 6.3: SHA promotion path on the earlier **n2-standard-4** multi-workload grid. Most configurations receive only cheap probes, while the active set shrinks as the path budget increases.

The progression illustrates why SHA is useful in this setting. Most configurations receive only cheap measurements, while promising configurations receive more budget. In the final experiment this is enough to identify the same best runtime configuration as the full grid: **rust/european/none** at $M = 100,000$, with runtime 0.54 ms.

6.14 Cross-Question Discussion

6.14.1 Software Stack and Engine Performance

The local-volatility results show that engine choice has a substantial effect on throughput. Rust/Rayon is the fastest no-AD engine by runtime across the main M -sweep and the large stress test. At $M = 500k$, $N = 252$, Rust is about $2.1\times$ faster than NumPy and about $1.25\times$ faster than JAX or C++. JAX and C++ are very similar at larger local-volatility sizes, while NumPy remains consistently slower.

The result also shows why the European GBM workload should not be treated as the main benchmark. On the small GBM baseline, C++ was clearly dominant. On local volatility, Rust becomes fastest and JAX becomes competitive with C++. This demonstrates that engine rankings depend strongly on workload structure.

The numerical interpretation is more cautious. The independent uncertainty check indicates that NumPy, C++ and JAX agree within estimated Monte Carlo error on the local-volatility price, while Rust does not. Thus Rust is best interpreted here as evidence for the throughput available from a Rayon-based native implementation, not as a fully validated local-volatility pricing backend.

6.14.2 Programming-Language, Runtime and AD-Support Trade-offs

The compiled CPU engines provide the strongest no-AD local-volatility runtime performance. Rust/Rayon performs best overall in the local-volatility timing experiments, while C++/OpenMP remains competitive and scales well with threads. JAX/XLA is not the fastest no-AD engine, but it provides the main AD capability and becomes close to C++ on larger multi-step workloads.

This should not be read as a broad benchmark of AD libraries across languages. The completed implementation compares JAX forward- and reverse-mode AD against native no-AD engines in NumPy, C++ and Rust. The C++ and Rust engines are valuable systems baselines, but they do not evaluate C++ AD tools, Rust AD libraries, Enzyme, or source-transformation AD in those languages.

The AD results show a clear trade-off. Reverse-mode AD is consistently much faster than forward-mode AD for the scalar-output local-volatility pricing function, with reverse-mode overhead usually around $2.1\text{--}2.6\times$ and forward-mode overhead around $4\text{--}7\times$. However, reverse mode also reaches memory limits at larger M, N combinations. Therefore, the best AD choice depends on whether the bottleneck is runtime or memory.

6.14.3 Parallelism and Bottlenecks

The thread-scaling results show that C++ and Rust expose path-level parallelism effectively to OpenMP and Rayon. Rust achieves $3.66\times$ strong-scaling speedup at four requested threads and $3.68\times$ weak-scaling throughput speedup. C++ also scales well, though less strongly, with $3.09\times$ strong-scaling speedup and $3.11\times$ weak-scaling throughput speedup.

JAX behaves differently. It is competitive in absolute runtime for no-AD local-volatility execution, but it does not respond to requested thread-count changes. This indicates that XLA’s managed CPU runtime should be interpreted separately from explicit OpenMP/Rayon thread scaling. For this project, the main bottlenecks are therefore different by engine: compiled engines are limited by parallel efficiency and memory bandwidth at higher thread counts, while JAX is limited by its managed execution model and, for reverse-mode AD, by memory pressure.

6.14.4 Compiler and AD Runtime Techniques

The JAX/XLA results provide evidence for the relevance of compiler and AD runtime techniques to Monte Carlo simulation. JAX is relatively weak on the small European GBM baseline, but becomes competitive with C++ on the multi-step local-volatility workload. This suggests that XLA-style compiled execution becomes more attractive when the workload has enough repeated structure to amortise dispatch and compilation overheads. The conclusion is therefore limited but useful: in this project, JAX/XLA is most valuable as an AD-capable compiled backend rather than as a universally fastest no-AD engine.

6.14.5 Integrated Scope of the Completed Evaluation

The completed experiments strongly address local MC+AD benchmarking, programming-language comparison, AD overhead, memory/resource limits, thread scalability, and a targeted form of cloud resource selection. The final cloud-profiler study adds measured cost-performance, Pareto-frontier, and profiler-selection evidence.

This scope is scientifically useful. The project demonstrates a reproducible benchmarking framework that can run multiple engines under matched workload definitions, capture repeated measurements, evaluate AD overheads, reveal practical resource limits, and attach cloud cost metadata to final selection decisions. The local-volatility suite also shows that more realistic MC+AD workloads can change the conclusions drawn from simple one-step baselines.

6.15 Limitations and Threats to Validity

6.15.1 Hardware and Cloud Scope

The detailed scaling and memory experiments were collected on a single Codespace environment. This makes the comparisons controlled and reproducible under the recorded conditions, but it prevents broad claims about AMD versus Intel, workstation versus cloud, or CPU versus GPU performance. The final cloud experiments add targeted Google Cloud evidence, but those results are still tied to specific instance types, regions, prices, software versions, and run dates. They should therefore be interpreted as measured cost-performance evidence for the reported configurations, not as universal cloud hardware rankings.

6.15.2 Local-Volatility Correctness

The European GBM workload has a closed-form Black–Scholes reference, but the chosen parametric local-volatility workload does not. Local-volatility correctness is therefore assessed through cross-engine consistency and convergence with larger M , rather than by exact analytical error. The additional high- M uncertainty check quantifies this limitation: NumPy, C++ and JAX agree within estimated Monte Carlo error, while Rust does not. For that reason, Rust local-volatility results are retained as throughput and scaling evidence, but Rust is not treated as a correctness-certified local-volatility pricing backend in the guarded-profiler deployment recommendation. Price differences between engines may reflect different random-number generators, path-generation implementations, or, in the Rust case, a backend-specific issue that requires further investigation. Likely causes include RNG stream partitioning, per-thread path assignment, discretisation mismatch, or a payoff/discounting implementation difference. A future validation step should use deterministic per-path random streams shared across engines and compare pathwise terminal values before comparing aggregate prices.

The AD correctness story is similarly scoped. European GBM prices and Greeks are analytically validated against Black–Scholes. Local-volatility prices are validated by cross-engine agreement and convergence where possible, but local-volatility Greeks have no closed-form reference in this report. The reported local-volatility sensitivities are therefore checked as internally consistent JAX AD outputs, not as analytically validated Greeks. A finite-difference check using common random numbers would be the most useful next validation step for the reported JAX local-volatility Greeks.

6.15.3 Random Number Generation as a Confounder

The engine comparisons measure complete backend implementations rather than isolated path-evolution kernels. This is useful for deployment because random number generation is part of real Monte Carlo runtime, but it also means the measured differences between NumPy, JAX, C++ and Rust combine simulation-kernel cost, RNG cost, memory layout, and reduction behaviour. The current results do not separate these effects. A stronger microbenchmark would run a pre-generated-normal experiment in which all engines consume the same normal variates, allowing RNG generation to be separated from path evolution and payoff reduction.

6.15.4 Warm versus Cold JAX Execution

The main timing tables report warm steady-state runtimes after configured warmup repetitions. This is appropriate for repeated production-style simulation and for comparing amortised kernel execution, but it is optimistic for short-lived cloud jobs where JAX/XLA compilation may be paid on the critical path. The implementation records and discusses compilation effects, but the final cloud cost-performance tables do not include a full cold-start versus warm-start comparison. Therefore, cloud cost claims should be read as steady-state cost-performance rather than one-off job latency.

6.15.5 Cloud Cost Scope

The cloud-cost estimates use marginal compute cost only. They intentionally exclude startup and shutdown time, idle time, persistent storage, orchestration overhead, retries, network transfer and billing granularity. This is defensible for comparing measured configurations under the same benchmark protocol, but it should not be interpreted as an exact Google Cloud invoice or as a complete end-to-end deployment-cost model.

6.15.6 Memory Measurement

Process-level memory measurements are coarse, especially for JAX/XLA, which may cache allocations or materialise intermediates inside compiled execution. For this reason, small RSS deltas are not over-interpreted. The OOM events are treated as stronger evidence than individual low memory-delta readings.

6.15.7 Thread Control

Thread-count control is more direct for C++/OpenMP and Rust/Rayon than for JAX/XLA. JAX uses its own CPU execution runtime, so requested OpenMP-style thread counts do not translate into clean JAX scaling experiments. This is why JAX thread results are interpreted as runtime behaviour rather than as a controlled thread-count benchmark.

6.16 Summary

Overall, the results show that MC+AD performance depends on workload structure, runtime system, differentiation mode, parallel execution model and cloud cost assumptions. European GBM validates prices and JAX Greeks against Black-Scholes, while the local-volatility suite exposes the main runtime, AD, memory and scaling trade-offs. Rust/Rayon gives the fastest measured no-AD local-volatility timings, but its local-volatility price discrepancy means those rows are performance evidence rather than correctness evidence. JAX provides the main AD evidence: reverse mode is faster than forward mode for moderate scalar-output local-volatility workloads, but reaches memory limits at larger M, N combinations.

The cloud results show that marginal cost-performance can differ from raw runtime ranking. Plain SHA halves full-grid work on the favourable **n2-standard-4** slice, while the guarded local-volatility profiler matches the observed oracle across the five-instance grid and avoids 13 of 30 full-scale task-level runs. These profiler results are promising, but scoped to the measured grids and should be read together with the limitations above.

Chapter 7

Conclusion

7.1 Summary of Work

This project designed, implemented, and evaluated a reproducible benchmarking framework for Monte Carlo option-pricing engines with automatic differentiation, with JAX providing the main implemented AD backend. The framework supports explicit workload configurations, engine abstraction, deterministic seeding, warm-up and repeated timing, analytical validation where available, SQLite result storage, environment metadata capture, cloud cost metadata, and generated report artefacts.

The completed implementation compares NumPy, JAX/XLA, C++/OpenMP, and Rust/Rayon engines across European GBM, parametric local-volatility, and Asian option workloads where supported. The local experimental campaign provides detailed evidence on validation, runtime, throughput, AD overhead, path-count scaling, time-step scaling, thread scaling, memory pressure, surface-shape sensitivity, and timing stability. The final cloud experiment extends this with cost-performance analysis, Pareto-frontier analysis, and profiler-versus-grid evaluation.

7.2 Key Findings

The results show that MC+AD performance cannot be reduced to a single backend ranking. The simple European GBM workload is useful for correctness checks, but the more demanding local-volatility and Asian workloads change the observed trade-offs. Rust/Rayon is the strongest no-AD engine by measured runtime in the main local-volatility results and in the highlighted European/local-volatility cloud rows, although the Rust local-volatility price requires further RNG-stream validation before it should be treated as correctness evidence. JAX/XLA remains essential for AD and is competitive for some path-dependent configurations.

The AD results show a clear runtime-memory trade-off. Reverse-mode AD is substantially cheaper than forward mode for the scalar-output local-volatility workload at moderate sizes, but it reaches practical memory limits at larger M, N combinations. In the final cloud comparison at $M = 100,000$, JAX AD overhead ranges from $1.81\times$ for Asian reverse mode to $4.42\times$ for local-volatility forward mode.

The cloud-profiler result is a useful final systems contribution. The earlier `sha_cloud_profiler_v4` run showed that plain SHA can halve the number of full benchmark cells on the favourable `n2-standard-4` slice, but also showed that unguarded pruning is too brittle across instance slices. The final guarded local-volatility experiment addresses this directly: across five four-vCPU Google Cloud instance types and 120 stored observations, guarded SHA matches the observed full-grid oracle for both the price-only and AD-required tasks. In the all-implemented-backend pool, the best observed price-only performance candidate is `rust/none/t2d-standard-4`; however, because Rust local-volatility pricing has not passed the same cross-engine consistency check as NumPy, C++ and JAX, the correctness-certified price-only recommendation is `cpp/none/t2d-standard-4`. When Greeks are required, the guarded profiler recommends `jax/reverse/t2d-standard-4`. These recommendations are scoped to the tested local-volatility grid and $100k$ -path target budget, where they achieve 0.0% objective regret while avoiding 13 of 30 full-scale task-level runs.

7.3 Answering Research Questions

RQ1. How do different software stacks and execution backends affect the runtime, throughput, and memory behaviour of Monte Carlo workloads? Engine choice materially changes runtime, throughput and memory behaviour. NumPy is a clear baseline, JAX/XLA benefits when workloads can be staged and compiled, and C++/OpenMP and Rust/Rayon provide strong native no-AD CPU throughput. JAX reverse-mode AD is the clearest memory-pressure case.

RQ2. What are the performance and memory trade-offs of automatic differentiation for representative Monte Carlo option-pricing workloads? AD is evaluated primarily through JAX. Reverse mode is usually faster than forward mode for scalar-output local-volatility pricing, but it has higher memory pressure and fails at larger path-count/time-step combinations. The comparison is therefore JAX AD versus native no-AD baselines, not a complete C++/Rust AD-library benchmark.

RQ3. How do workload structure and scale, including path count, time-step count, and payoff type, affect backend rankings and bottlenecks? Rankings change between European GBM, local volatility and Asian payoffs. Path count mainly tests parallel throughput; local-volatility time steps expose scan cost and AD memory growth. Rust local-volatility timings are among the fastest measured timings, subject to the Chapter 6 correctness caveat.

RQ4. Which Google Cloud CPU configurations provide the best marginal cost-performance for the evaluated MC+AD workloads? Marginal cost-performance is not the same as raw speed. In the final guarded local-volatility run, `t2d-standard-4` wins the weighted objectives because it is close to the fastest `c2d-standard-4` timings while being cheaper per measured run. These are marginal compute estimates, not complete cloud bills.

RQ5. Can short probe runs and successive-halving style profiling reduce the cost of cloud benchmarking while still recovering near-Pareto-optimal configurations? Yes within the measured grids, but not universally. Plain SHA recovered the full-grid Pareto frontier in the main `n2-standard-4` slice while halving full-scale evaluations, but cross-instance robustness was weaker. Guarded SHA recovered the observed local-volatility oracle while avoiding 13 of 30 full-scale task-level runs.

RQ6. What are the practical limitations, threats to validity, and generalisation limits of this benchmarking methodology? The key limits are finite candidate grids, warm rather than full cold-start JAX timings, marginal rather than end-to-end cloud costs, different RNG implementations, limited GPU evidence, and no closed-form local-volatility reference. The results support measured systems claims, not universal rankings of languages, AD systems or cloud hardware.

7.4 Future Work

The most direct next step is to validate the guarded cloud-profiler policy with more workloads, more instance types, repeated cloud runs, current pricing snapshots, and uncertainty estimates for timing variance, scaling crossovers, and Monte Carlo noise. Local-volatility Greek validation should also be strengthened with finite-difference checks using common random numbers. Other useful extensions are GPU and distributed execution, mixed precision, hardware counters, and broader AD-system coverage such as Enzyme, C++ AD tools, Rust AD libraries, Mojo, or specialist quant libraries. Two focused microbenchmarks would also sharpen the performance interpretation: pre-generated normals to separate RNG cost from path evolution, and cold-start versus warm-start JAX timing to decide when XLA compilation should be included in cloud cost-performance decisions. The frontend/API could then be extended into a fuller interactive dashboard for profiler configuration and report-artefact generation.

7.5 Final Remarks

Overall, the project contributes both a working engineering artefact and an empirical methodology. The artefact makes MC+AD experiments reproducible, queryable, and extensible across engines. The methodology shows how local benchmarking can be combined with cloud cost metadata and successive halving to reduce full-grid benchmarking, while the guarded-SHA experiment shows how explicit safeguards can make that reduction more defensible. The profiler result is therefore promising evidence for budget-aware benchmarking, not a universal guarantee of optimal cloud selection.

Declarations

Academic Integrity

This report is submitted as my own work. External sources, software libraries, and project guidance have been acknowledged where used. The implementation, experimental design, analysis, and final interpretation were reviewed and assembled by the author.

Generative AI Assistance

Generative AI tools were used as writing and coding assistants during parts of the project, including for drafting support, proofreading, code explanation, and editing suggestions. All AI-assisted material was critically reviewed, checked against the project repository, benchmark artefacts, and generated tables or figures, and edited by the author before inclusion. Numerical claims in the report are based on stored benchmark results and generated report artefacts rather than unsupported model output. The final interpretation, selection of evidence, and responsibility for the submitted work remain mine.

Bibliography

- [1] Paul Glasserman. *Monte Carlo Methods in Financial Engineering*. Springer, New York, 2004.
- [2] Andreas Griewank and Andrea Walther. *Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation*. SIAM, Philadelphia, 2 edition, 2008.
- [3] Atilim Gunes Baydin, Barak A. Pearlmutter, Alexey Andreyevich Radul, and Jeffrey Mark Siskind. Automatic differentiation in machine learning: A survey. *Journal of Machine Learning Research*, 18(153):1–43, 2018.
- [4] James Bradbury, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, and Skye Wanderman-Milne. JAX: Composable transformations of Python+NumPy programs. <https://github.com/jax-ml/jax>, 2018.
- [5] William S. Moses and Valentin Churavy. Enzyme: High-performance automatic differentiation of llvm. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, 2021.
- [6] Michael Innes. Zygote: A differentiable programming system for julia. *arXiv preprint arXiv:1907.07587*, 2019.
- [7] Yuanming Hu, Shuchang Liu, Matthias Zwicker, Zhendong Su, and Yizhou Sun. DiffTaichi: Differentiable programming for physical simulation. *arXiv preprint arXiv:1910.00935*, 2020.
- [8] Michael B. Giles and Paul Glasserman. Smoking adjoints: Fast monte carlo greeks. *Risk*, 19(1):88–92, 2006.
- [9] Luca Capriotti. Fast greeks by algorithmic differentiation. *Journal of Computational Finance*, 14(3):3–35, 2011.
- [10] Lukas Weber. Carlo.jl: A general framework for monte carlo simulations in julia. *SciPost Physics Codebases*, 49, 2025.
- [11] Louis-Noël Pouchet. Polybench/c 3.2, 2012. Benchmark suite.
- [12] Jan Treibig, Georg Hager, and Gerhard Wellein. Likwid: A lightweight performance-oriented tool suite for x86 multicore environments. In *2010 39th International Conference on Parallel Processing Workshops*, 2010.
- [13] Intel Corporation. Intel vtune profiler, 2025. Software documentation.
- [14] OpenXLA Project. Xprof: Profiling for xla / openxla, 2025. Project documentation.
- [15] Chris Lattner et al. Mlir: A compiler infrastructure for the end of moore’s law. In *Proceedings of the 3rd ACM SIGPLAN International Workshop on Machine Learning and Programming Languages*, 2020.
- [16] Tianqi Chen, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Haichen Shen, et al. Tvm: An automated end-to-end optimizing compiler for deep learning. In *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation*, 2018.
- [17] Lianmin Zheng, Chengfan Jia, Minmin Sun, Zhao Wu, Cody Hao Yu, Ahmad Haj-Ali, et al. Anso: Generating high-performance tensor programs for deep learning. In *Proceedings of the 14th USENIX Symposium on Operating Systems Design and Implementation*, 2020.

- [18] Paulius Micikevicius, Sharan Narang, Jonah Alben, Gregory Diamos, Erich Elsen, David Garcia, et al. Mixed precision training. *arXiv preprint arXiv:1710.03740*, 2017.
- [19] Cody Coleman, Deepak Narayanan, Daniel Kang, Tian Zhao, Jian Zhang, Luigi Nardi, et al. Dawnbench: An end-to-end deep learning benchmark and competition. *arXiv preprint arXiv:1706.04567*, 2017.
- [20] Cody Coleman, Deepak Narayanan, Daniel Kang, Tian Zhao, Jian Zhang, Luigi Nardi, et al. An analysis of deep learning training times with dawnbench. *arXiv preprint arXiv:1908.00799*, 2019.
- [21] Peter Mattson, Christine Cheng, Greg Diamos, Cody Coleman, Paulius Micikevicius, David Patterson, et al. Mlperf training benchmark. *Proceedings of Machine Learning and Systems*, 2, 2020.
- [22] Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973.